

Probability Theory

Jacob Denson

June 7, 2020

Table Of Contents

1	Foundations	2
1.1	Frequentist Probability	3
1.2	Bayesian Probability	4
1.3	Axioms of Probability	7
1.4	Conditional Probabilities	13
1.5	Kolmogorov's Zero-One Law	16
1.6	Beware of Intuition	16
2	Random Variables	18
2.1	Distributions	19
2.2	Expectation	19
3	Tail Bounds and Concentration	20
3.1	Tail Bounds	22
3.2	Moments of Sums of Random Variables	24
3.3	Concentration of Measure	28
3.4	Subgaussian Random Variables	31
3.5	Subexponential Random Variables	37
4	Existence Theorems	39
5	Entropy	40
6	Appendix: Uniform Integrability	43
7	Percolation Theory	48
7.1	Duality	49
7.2	Boolean Functions and Sharp Thresholds	49

7.3	Conformal Invariance	52
8	High Dimensional Probability	54

Chapter 1

Foundations

These notes outline the basics of probability theory, the mathematical framework which allows us to interpret the statement that we are 8 times more likely to develop lung disease if you are smoker than if you are a non-smoker, or that there is a 50% *chance* of rain on Saturday?

These statements seem intuitive, and we use them naturally in everyday conversation. But a closer analysis of these statements reveals a couple difficulties with understanding such a statement. On Saturday, it will either rain, or not rain, so it is nontrivial to see how one would calculate a reasonable universal ‘chance’ of such an event happening.

Mathematicians have a rigorous abstraction of these statements, interpreted in the language of measure theory. Soon, we will begin to work in this language. But probability theory is also firmly grounded in scientific applications. So at least for intuition, we should begin by exploring the various interpretations of probability theory in real life applications.

In this chapter, we will explore the two major interpretations of probability theory in real life, each of which use the same underlying mathematical theory to make judgements about the world. But regardless of which interpretation you have, the axiomatic system through which mathematicians study probability remains the same; the two interpretations differ only through which the system is applied to model real life events.

1.1 Frequentist Probability

Classical probability theory was developed according to the intuitions of what is now known as the frequentist school of probability theory, and is the simplest interpretation of probability to understand. It is most easily understood from the point of view of a scientist. Suppose you are repeatedly performing some well-controlled experiment, in the sense that you expect a similar outcome to occur in each trial. Even under rigorously controlled conditions, the experiment will not always result in the same outcome. Slight experimental error results in slight changes in the outcome of the experiment. Nonetheless, some outcomes may occur more frequently than others.

Let us perform an experiment as often as desired, obtaining an infinite sequence of results. Let E be a certain statement about the outcome of the experiment. In probability theory, we call this an **event**. For instance, E may ask whether a flipped coin lands heads up when flipping a coin repeatedly. We define the **relative frequency** of E being true in n trials by the equation

$$P_n(E) := \frac{\#\{k \leq n : E \text{ is true for experiment } k\}}{n}$$

The key assumption of the frequentist school of probability is the existence of a *long term relative frequency* for these events; if our experiments are suitably controlled, then regardless of the specific sequence of measured outcomes, our relative frequencies will always converge to a well defined invariant ratio. We define this ratio to be the **probability** of a certain event, denoted $\mathbf{P}(E)$. That is,

$$\mathbf{P}(E) := \lim_{n \rightarrow \infty} P_n(E).$$

Let's explore some consequences of this doctrine:

- Since $0 \leq P_n(E) \leq 1$ for any n , taking limits shows $0 \leq \mathbf{P}(E) \leq 1$.
- If Ω is a tautological statement, i.e. an event which is true for any experiment, then $P_n(E) = 1$ for all n , so

$$\mathbf{P}(\Omega) = 1.$$

- If E and F are disjoint propositions, in the sense that both outcomes can never be simultaneously true, then

$$\begin{aligned} P_n(E \vee F) &= \frac{\#\{k \leq n : E \vee F \text{ is true for experiment } k\}}{n} \\ &= \sum_{i=1}^{\infty} \frac{\#\{k \leq n : E_i \text{ is true for experiment } k\}}{n} = \sum_{i=1}^{\infty} P_n(E_i). \end{aligned}$$

Thus

$$\mathbf{P}\left(\bigvee_{i=1}^{\infty} E_i\right) = \sum \mathbf{P}(E_i).$$

The properties described here turn out to be sufficient to describe all the mathematically important rules of probability theory. What's more, we can use these rules as axioms to *prove* that the relative frequencies of a sequence of controlled experiments eventually settles down to a well defined ratio, a fact known as the strong and weak laws of probability, which justifies the thought process of the frequentist school in the first place.

1.2 Bayesian Probability

The frequentist school is sufficient to use probability theory to model scientific experiments, but in everyday life we make a more expansive use of probabilistic language. If you turn on the news, it's common to hear that "there is an 80% chance of downpour this evening". It is difficult to interpret this result in the frequentist definition of probability. Even if we see each night's temperament as an experimental trial, it is hard to convince yourself that these experiments are controlled enough to converge to a probabilistic result. The Bayesian school of probability redefines probability theory to be attuned to a person's individual beliefs, so that we can interpret "there is an 80% chance of downpour this evening" as an individual's belief that they think it will rain this evening rather than not rain.

You might argue that, if probability is a personal belief in an unknown event, we can choose probabilities however we want, which would break down the logical structure required to make any mathematical progress. However, the probabilities that the Bayesian school studies are forced to be 'logically consistent'. This forces the logical structure required for mathematical probability theory.

Consistency can be formulated in various ways; here we discuss what is known as the Dutch book method, developed by the Italian probabilist Bruno de Finetti; if you assign to a certain unknown event E a probability $\mathbf{P}(E) \in [0, 1]$, then you believe that a bet at $[A \cdot \mathbf{P}(E) : A \cdot (1 - \mathbf{P}(E))]$ odds is completely fair for any $A > 0$. In other words, you would be willing to make a bet that if E occurs, you win $A \cdot (1 - \mathbf{P}(E))$ dollars, and if E does not occur, you have to pay up $A \cdot \mathbf{P}(E)$ dollars. Since you believe the bet is fair, you must also be willing to play a game where you lose $A \cdot (1 - \mathbf{P}(E))$ dollars if E occurs, and gain $A \cdot \mathbf{P}(E)$ dollars if E does not occur. For instance, you might be willing to bet a dollar against a dollar that a coin will turn up heads, which is $[1 : 1]$ odds, so we would assign the probability that a coin will turn up heads as $1/2$.

We saw a probability function \mathbf{P} is **inconsistent** if it possible to make a series of bets that will guarantee a profit regardless of the outcome; this is known as a Dutch book. We view such functions as illogical, and so in probability theory we concentrate solely on **consistent** probability functions. It turns out that consistent probability functions satisfy the same properties as frequentist probabilities:

- Suppose an event Ω is tautological, and thus always occurs. We claim that any consistent probability function \mathbf{P} must satisfy

$$\mathbf{P}(\Omega) = 1.$$

Suppose $\mathbf{P}(\Omega) < 1$. If we made a 1 dollar bet on Ω occurring, then we would make $1 - \mathbf{P}(\Omega) > 0$ dollars off the bet. This bet always makes money, so it is a Dutch book, and therefore \mathbf{P} is inconsistent.

- Suppose E_1, \dots, E_n are disjoint propositions. We claim that if \mathbf{P} is consistent, then

$$\mathbf{P}\left(\bigvee_{i=1}^n E_i\right) = \sum_{i=1}^n \mathbf{P}(E_i).$$

Suppose \mathbf{P} is a probability function such that

$$\mathbf{P}\left(\bigvee_{i=1}^n E_i\right) < \sum_{i=1}^n \mathbf{P}(E_i).$$

Make a $[\mathbf{P}(E_i) : 1 - \mathbf{P}(E_i)]$ bet that each event E_i does *not* occur, for each index i , as well as a $[\mathbf{P}(\bigvee_{i=1}^n E_i) : 1 - \mathbf{P}(\bigvee_{i=1}^n E_i)]$ bet that

$\bigvee_{i=1}^n E_i$ *does* occur. Suppose the event E_i occurs. Then we win every bet that E_j does not occur, for $j \neq i$, as well as the event that $\bigvee_{i=1}^n E_i$ occurs. Our profits are

$$\left(1 - \mathbf{P}\left(\bigvee_{i=1}^n E_i\right)\right) + \sum_{j \neq i} \mathbf{P}(E_j)$$

and we lose $1 - \mathbf{P}(E_i)$ dollars. Thus our total revenue is

$$\begin{aligned} \left(1 - \mathbf{P}\left(\bigvee_{i=1}^n E_i\right)\right) + \sum_{j \neq i} \mathbf{P}(E_j) - (1 - \mathbf{P}(E_i)) \\ = \sum \mathbf{P}(E_i) - \mathbf{P}\left(\bigvee_{i=1}^n E_i\right) > 0. \end{aligned}$$

On the other hand, if $\bigvee_{i=1}^n E_i$ does *not* occur, then we win every bet that E_i does *not* occur, but lose the bet that $\bigvee_{i=1}^n E_i$ occurs. Thus our profits are $\sum \mathbf{P}(E_i)$ dollars, but have to pay up $\mathbf{P}(\bigvee_{i=1}^n E_i)$ dollars. Our total revenue is

$$\sum \mathbf{P}(E_i) - \mathbf{P}\left(\bigvee_{i=1}^n E_i\right) > 0.$$

Thus we always make a positive revenue, so we have a Dutch book. On the other hand, if

$$\mathbf{P}\left(\bigvee_{i=1}^n E_i\right) > \sum_{i=1}^n \mathbf{P}(E_i),$$

then making the opposite series of bets also gives a Dutch book. Thus both inequalities gives a contradiction, so the required identity holds for any consistent probability function.

One technical difference in the Bayesian regime is that we are only able to prove that $\mathbf{P}(\bigvee E_i) = \sum \mathbf{P}(E_i)$ for *finite* disjoint unions of events. But allowing countable sums is so useful that it is difficult to use finite unions (though Defineti forced himself only to finite unions). But if you are willing to let this slight technical problem slip by, this means that Bayesian

and Frequential probability theories both lead to the same fundamental laws of probability. We shall take the two laws we have derived, and use them to make a rigorous model so no more philosophical questions can enter the theory. This is where mathematical probability theory takes its form.

1.3 Axioms of Probability

Probability theory is rigorously defined using measure theory. We consider a set Ω , which we view as a set of ‘all possible outcomes’ to some random phenomenon. This is called the **sample space**. We then fix a sigma-algebra Σ on Ω , whose elements we call **events**, and consider a positive measure $\mathbf{P} : \Sigma \rightarrow [0, \infty)$ with $\mathbf{P}(\Omega) = 1$, known as a **probability measure**. Together, $(\Omega, \Sigma, \mathbf{P})$ gives the structure of a **probability space**. We view Σ as the set of all ‘acceptable’ statements to ask the probability of. Measure theory tells us that we cannot assign probabilities to all subsets of Ω for sufficiently complicated probability distributions, which is the reason for this technicality.

From this summary, it seems probability theory should be viewed as a subset of measure theory, studying positive measures. But probability theory really only takes a probability space as the *model* on which to study certain phenomena. Once this is done, probability takes a completely different viewpoint to measure theory, which leads to completely different mathematical questions being raised. For instance, in probability theory we hardly ever concentrate on more than a single probability space, which we view as encapsulating all the randomness in a particular situation. Furthermore, as we delve deeper into the theory the technical aspects of measure theory tend to disappear, as the probabilistic mindset becomes much more natural.

Classically, probability theory was the study of certain techniques used to calculate $\mathbf{P}(E)$ for certain outcomes $E \subset \Omega$, which motivated the development of the modern fields of combinatorics and integration theory. Nowadays, probability theory tends to focus more on general principles underlying probability spaces, and integrates tools in more modern fields such as functional analysis to achieve this.

Example. *Consider the situation of flipping a coin, and observing a heads or tails. Let’s model this random situation as a probability space. The sample*

space is $\Omega = \{H, T\}$, which corresponds to all outcomes of the coin flip. Since the sample space is suitably discrete, we can let the sigma algebra consist of all subsets of Ω , i.e. we set $\Sigma = 2^\Omega$. Since we view both heads and tails as being equally likely, we set $\mathbf{P}(H) = \mathbf{P}(T) = 1/2$. This extends to a probability measure on Σ by setting, for each $S \subset \Omega$,

$$\mathbf{P}(S) = \sum_{\omega \in S} \mathbf{P}(\omega)$$

This is an atomic way of defining a probability space, since we define the probability of certain atomic events occurring, and then extend the definition to all events by taking sums.

Example. Given any finite sample space Ω , we can put a **uniform distribution** on Ω atomically, by defining $\mathbf{P}(s) = 1/|\Omega|$ for each $s \in \Omega$.

Example. If $\omega \in \Omega$ is fixed, we can put a **point mass distribution** on Ω by defining, for each $E \subset \Omega$,

$$\mathbf{P}(E) = \begin{cases} 1 & \omega \in E \\ 0 & \omega \notin E \end{cases}$$

The distribution represents an event where a single outcome is certain to occur, and all other situations are impossible.

There is no restriction in mathematical probability theory on *how* particular probabilistic events are obtained, or what probabilities are assigned to each event. All we need is that the model studied agrees with the laws of probability. For instance, in classical statistical mechanics, one can often use a discrete model consisting of placing a certain number of indistinguishable molecules in a certain number of positions. If the positions are indistinguishable, it is reasonable to assume that each molecule independently occurs in any position with equal probability, an assumption leading to the Maxwell-Boltzmann theory of statistical mechanics. Given n points to place in m positions, combinatorics tells us the probability that k_1 points occur in the first position, k_2 in the second, and so on up to k_m in the m 'th position, is

$$\frac{1}{m^n} \binom{n}{k_1 \ k_2 \ \dots \ k_m} = \frac{1}{m^n} \frac{n!}{k_1! k_2! \dots k_m!}$$

Thus more evenly spread states are more likely to occur. However, in the 20th century Bose and Einstein found that in the study of certain particles, any such configuration (k_1, \dots, k_m) has an *equal* chance of occurring, leading to a completely different assignment of probabilities. To the mathematician, both theories are an equally applicable area of study.

Example. If Ω is a countable set, then a probability measure on Ω as a member of the set

$$\left\{ \mathbf{P} : \Omega \rightarrow [0, 1] : \sum_{\omega \in \Omega} \mathbf{P}(\omega) = 1 \right\} \subset [0, 1]^\Omega$$

This is a convex subset of the unit ball in $l^\infty(\Omega)$, which leads to some interesting linear analysis.

When studying countable probability spaces, the sigma algebra plays almost no real role in the theory. This allows us to get away with discussing most of the basic principles of probability theory without running into too many technicalities. Nonetheless, even in the study of discrete phenomena understanding probability spaces with uncountably many points becomes necessary. For instance, in the study of the limiting average of a sequence of discrete coins flips, our sample space must consist of the space of infinite sequences of coin flips $\{0, 1\}^\omega$, which is an uncountable sample space. And it is often necessary in applications to select a point in an interval uniformly at random.

Example. Consider a five digit number X selected uniformly at random. Then the sample space consists of the numbers $[0, 99999]$, and the probability that a particular number is selected is one in a 100,000. There are $10!/5! = 30240$ numbers all of whose digits are different, so the probability that a number is selected all of whose digits are different is 0.3024. If we look at the first 800 digits of the decimal expansion of the number e , and we take each 5 digit consecutive sequence in this expansion, we end up with approximately the same frequency of unique digits, leading us to believe the digits occurring in the number e are essentially random.

Example. An interesting application of combinatorial probability theory is in the so called Birthday paradox. Given a number of n points to place uniformly in m boxes, the probability that they all lie in distinct boxes is

$$\frac{1}{m^n} \frac{m!}{(m-n)!} = \prod_{k=0}^{n-1} (1 - k/m).$$

In particular, if $m = 365$, and $n = 23$, then we calculate the probability that two points lie in the same box exceeds one half. To determine this approximately, we take logarithms, using the fact that $\log(1 - x) = -x + O(x^2)$ for small values of x , so

$$\log \left(\prod_{k=1}^n (1 - k/m) \right) = - \sum_{k=1}^n k/m + O(k^2/m^2) = -\frac{n(n+1)}{2m} + O(n^3/m^2).$$

Thus the probability that there exist two points lying in the same box

$$1 - \prod_{k=1}^n (1 - k/m) = 1 - \exp \left(-\frac{n(n+1)}{2m} + O(n^3/m^2) \right) = \frac{n(n+1)}{2m} + O(n^4/m^2).$$

Thus if $n \gtrsim m^{1/2}$, it is highly likely for two points to lie in the same box. In particular, if $m = 365$, then this estimate says that in groups of more than 27 people, there is more than 50% chance that two people share the same birthday.

The first immediately obvious fact from the axioms is $\mathbf{P}(E^c) = 1 - \mathbf{P}(E)$, since E and E^c are disjoint events whose union is Ω . A similar discussion shows that $\mathbf{P}(E \cup F) = \mathbf{P}(E) + \mathbf{P}(F) - \mathbf{P}(E \cap F)$ because $E \cup F$ can be written as the union of the three disjoint events $E \cap F$, $E \cap F^c$, and $E^c \cap F$, and

$$\mathbf{P}(E) = \mathbf{P}(E \cap F) + \mathbf{P}(E \cap F^c) \quad \mathbf{P}(F) = \mathbf{P}(E \cap F) + \mathbf{P}(E^c \cap F)$$

This process can be generalized to unions of finitely many events. We have

$$\mathbf{P}(E \cup F \cup G) = \mathbf{P}(E) + \mathbf{P}(F) + \mathbf{P}(G) - \mathbf{P}(E \cap F) - \mathbf{P}(E \cap G) - \mathbf{P}(F \cap G) + \mathbf{P}(E \cap F \cap G)$$

which can be reasoned by looking at the number of times each element of $E \cup F \cup G$ is ‘counted’ on the right hand side. In general, we have the inclusion-exclusion principle

$$\mathbf{P} \left(\bigcup_{k=1}^n E_k \right) = \sum_{S \subset \{1, \dots, n\}} (-1)^{|S|} \mathbf{P} \left(\bigcap_{k \in S} E_k \right)$$

This can be proven by a clumsy inductive calculation. More interestingly, but less useful, we often want to calculate the probability of an infinite

union of sets E_k occurring. The inclusion-exclusion principle can be taken ‘in the limit’ to conclude that

$$\mathbf{P}\left(\bigcup_{k=1}^{\infty} E_k\right) = \lim_{n \rightarrow \infty} \mathbf{P}\left(\bigcup_{k=1}^n E_k\right) = \sum_{\substack{S \subset \mathbf{N} \\ |S| < \infty}} (-1)^{|S|} \cdot \mathbf{P}\left(\bigcap_{k \in S} E_k\right)$$

where the sum on the right is taken as the limit of the partial sums where $S \subset \{1, \dots, n\}$ – the sum need not converge absolutely, so it is important to take the limit in the precise ordering given.

The inclusion-exclusion formula can be tricky to calculate in real examples, so we often rely on estimates to upper bound or lower the probability of a particular event occurring. The trivial **union bound**

$$\mathbf{P}\left(\bigcup E_i\right) \leq \sum \mathbf{P}(E_i)$$

can often be applied. This is a good inequality to apply if the E_i are ‘nearly disjoint’, or each have a negligible probability of occurring. On the other hand, the bound is shockingly bad if all the E_i are equal to one another.

Another useful fact to consider is that $\mathbf{P}(E_k) \rightarrow \mathbf{P}(E)$ if the sets E_k ‘tend to’ E in one form or another. If the E_k are an increasing sequence whose union is E , then we can certainly conclude $\mathbf{P}(E_k) \rightarrow \mathbf{P}(E)$. Similarly, if E_k is a decreasing sequence whose intersection is E , then $\mathbf{P}(E_k) \rightarrow \mathbf{P}(E)$. To obtain general results, we say that $E_k \rightarrow E$ if $\limsup E_k = \liminf E_k = E$, where

$$\begin{aligned} \limsup_{k \rightarrow \infty} E_k &= \bigcap_{n=1}^{\infty} \bigcup_{k \geq n} E_k = \{\omega : \omega \in E_k \text{ for infinitely many } k\} \\ \liminf_{k \rightarrow \infty} E_k &= \bigcup_{n=1}^{\infty} \bigcap_{k \geq n} E_k = \{\omega : \omega \in E_k \text{ for sufficiently large } k\} \end{aligned}$$

We can then conclude that $\mathbf{P}(E_k) \rightarrow \mathbf{P}(E)$, since once can show

$$\limsup \mathbf{P}(E_k) \leq \mathbf{P}(\limsup E_k)$$

$$\liminf \mathbf{P}(E_k) \geq \mathbf{P}(\liminf E_k)$$

so we can apply the squeeze theorem. This already enables us to prove a very interesting theorem which can guarantee an event can ‘never occur’.

Lemma 1.1 (Borel-Cantelli Lemma). *If E_1, E_2, \dots is a sequence of events with $\sum \mathbf{P}(E_k) < \infty$, then $\mathbf{P}(\limsup E_k) = 0$. Thus none of the events E_k can happen infinitely often.*

Proof. Because

$$\mathbf{P}\left(\bigcup_{k \geq n} E_k\right) \leq \sum_{k \geq n} \mathbf{P}(E_k)$$

for any $\varepsilon > 0$ we can find an N such that for $n \geq N$, $\mathbf{P}(\bigcup_{k \geq n} E_k) < \varepsilon$. But for any n , $\limsup E_k \subset \bigcup_{k \geq n} E_k$, and so we conclude $\mathbf{P}(\limsup E_k) < \varepsilon$. We then let $\varepsilon \rightarrow 0$ to conclude $\mathbf{P}(\limsup E_k) = 0$. \square

The next example shows that the hypothesis $\sum \mathbf{P}(E_k) < \infty$ cannot be relaxed without further analysis of the events E_k beyond their probabilities.

Example. *Take the Haar measure measure on $\mathbf{T} = \mathbf{R}/\mathbf{Z}$. Consider a sequence of positive numbers x_1, x_2, \dots , define $S_N = \sum_{n=1}^N x_n$, and $E_n = [S_{n-1}, S_n]$, considered modulo \mathbf{Z} of course. Then $\mathbf{P}(E_n) = x_n$, and $\sum x_n = \infty$ happens if and only if every point in \mathbf{T} is contained in infinitely many of the E_n .*

Theorem 1.2. *If E_1, E_2, \dots are events with $\inf \mathbf{P}(E_k) > 0$, then infinitely many of the E_i occur at once with positive probability.*

Proof. The event that infinitely many of the E_1, E_2, \dots occur is the complement of the event that all but finitely many of the E_i do not occur, i.e. $\liminf E_i^c$, and it suffices to show $\mathbf{P}(\liminf E_i^c) < 1$. But by Fatou's lemma,

$$\mathbf{P}\left(\inf_{k \geq n} E_k^c\right) \leq \inf_{k \geq n} \mathbf{P}(E_k^c) = 1 - \sup_{k \geq n} \mathbf{P}(E_k) \leq 1 - \delta$$

and so, letting $n \rightarrow \infty$, we conclude $\mathbf{P}(\liminf E_k^c) \leq 1 - \delta$. Alternatively, if we consider the functions $S_n = \chi_{E_1} + \dots + \chi_{E_n}$, then

$$S_n \leq m \mathbf{I}(S_n \leq m) + n \mathbf{I}(S_n > m) = m + (n - m) \mathbf{I}(S_n > m)$$

so if $\delta = \inf \mathbf{P}(E_i)$, then

$$\delta n \leq \mathbf{E}(S_n) \leq m + (n - m) \mathbf{P}(S_n > m)$$

which leads to the upper bound

$$\mathbf{P}(S_n > m) \geq \frac{\delta n - m}{n - m}$$

As $n \rightarrow \infty$, the events on the left hand side increasing to $\mathbf{P}(S_\infty > m)$, where we define S_∞ as the sum of all χ_{E_k} . Thus

$$\mathbf{P}(S_\infty > m) \geq \limsup_{n \rightarrow \infty} \frac{\delta n - m}{n - m} = \delta$$

But we can then let $m \rightarrow \infty$ to conclude that $\mathbf{P}(S_\infty = \infty) = \delta$. \square

1.4 Conditional Probabilities

In the Bayesian interpretation of probability theory, it is natural for probabilities to change over time as more information is gained about the system in question. That is, given that we know some proposition F holds over the sample space, we obtain a new probability measure over Ω , denoted $\mathbf{P}(\cdot|F)$, which represents the ratio of winnings from the bet which is only played out if F occurs. That is

- You win $1 - \mathbf{P}(E|F)$ dollars if E occurs, and F occurs.
- You lose $\mathbf{P}(E|F)$ dollars if E does not occur, and F occurs.
- No money exchanges hands if F does not occur.

Suppose we had to assign values to $\mathbf{P}(E|F)$, subject to a consistency principle which prevents a dutch book argument. It then follows that we must have $\mathbf{P}(F)\mathbf{P}(E|F) = \mathbf{P}(E \cap F)$ TODO: Fill in this argument.

In the empirical interpretation, $\mathbf{P}(E|F)$ is the ratio of times that E is true in experiments, where we only count experiments in which F also occurs. That is, we define $\mathbf{P}(E|F)$ as the limit of the ratios

$$P_n(E|F) = \frac{\#\{k \leq n : \omega_k \in E, \omega_k \in F\}}{\#\{k \leq n : \omega_k \in F\}}$$

But it is easy to calculate, by dividing the numerator and denominator by n , that $P_n(E|F) = P_n(E \cap F)/P_n(F)$, so by taking limits, we find

$$\mathbf{P}(E|F) = \lim_{n \rightarrow \infty} \frac{P_n(E \cap F)}{P_n(F)} = \frac{\mathbf{P}(E \cap F)}{\mathbf{P}(F)}$$

which gives us the formula $\mathbf{P}(F)\mathbf{P}(E|F) = \mathbf{P}(E \cap F)$. We must of course assume that $\mathbf{P}(F) \geq 0$, since otherwise we are almost certain that F will

never occur, and then we can almost guarantee that the limit of the values $P_n(E|F)$ does not exist.

Thus we have motivation to define conditional probabilities by the formula $\mathbf{P}(F)\mathbf{P}(E|F) = \mathbf{P}(E \cap F)$, provided that $\mathbf{P}(F) > 0$. It enables us to model the information gained by restricting our knowledge to a particular subset of sample space. In particular, we can use the definition to identify events which contain information ‘useless’ to learning about another event. We say two events E and F are independent if $\mathbf{P}(E \cap F) = \mathbf{P}(E)\mathbf{P}(F)$, or, provided $\mathbf{P}(F) > 0$, $\mathbf{P}(E|F) = \mathbf{P}(E)$; knowledge of F gives us no foothold over knowledge of the likelihood of E .

Example. *The Monty Hall problem is an incredible example of how paradoxical probability theory can seem. We are on a gameshow. Suppose there are three doors in front of you. A car (brand new!) is placed uniformly randomly behind one of the doors. After we pick a door (the first door, for instance), the gameshow host then opens the second door, which you didn’t pick, revealing the car isn’t behind the door. It is important to note that he picked randomly from the remaining doors which you didn’t pick and don’t have a car behind them. What is the chance that the door you picked has the brand new car? You likely would think the two doors have a 50-50 chance of containing the car given this info, but you’d be wrong. Let $X \in \{1, 2, 3\}$ denote the door chosen uniformly at random where the car lies, and let $Y \in \{1, 2, 3\}$ denote the door that the host randomly chose to open. We know $Y \neq 1$, because the gameshow host would never open the door we picked; that would give the game away! If $X = 1$, then Y is picked from $\{2, 3\}$ with uniform possibility. However, if $X = 2$, something interesting occurs – the gameshow is forced to open door number 3, because that’s the only door that (he thinks) won’t give any information to the player, and similarly, if $X = 3$, then $Y = 2$. Now we know that since X is chosen uniformly at random $\mathbf{P}(X = k) = 1/3$ for each k . Similarly, we know that Y is then chosen uniformly at random from $\{2, 3\}$, given that $X = 1$, so assuming X and Y are independent, we conclude*

$$\mathbf{P}(X = 1, Y = 2) = \mathbf{P}(X = 1)\mathbf{P}(Y = 2) = 1/6$$

$$\mathbf{P}(X = 1, Y = 3) = 1/6$$

But we also know that if $X = 2$, then $Y = 3$, so

$$\mathbf{P}(X = 2, Y = 3) = \mathbf{P}(X = 2) = 1/3$$

$$\mathbf{P}(X = 3, Y = 2) = \mathbf{P}(X = 3) = 1/3$$

It follows that

$$\begin{aligned} & \mathbf{P}(\text{door 1 has a car} | \text{door 2 was opened}) \\ &= \frac{\mathbf{P}(\text{door 1 has a car, door 2 was opened})}{\mathbf{P}(\text{door 2 was opened})} \\ &= \frac{\mathbf{P}(\{(X = 1, Y = 2)\})}{\mathbf{P}(\{(X = 1, Y = 2), (X = 3, Y = 2)\})} = \frac{1/6}{1/6 + 1/3} = 1/3 \end{aligned}$$

This means we should definitely change our minds about which door we were going to pick! The argument above causes a great media uproar when it was published in 1990 in a popular magazine, because of how convincing the fallacious argument below is. The total number of possibilities is

$$(X = 1, Y = 2), (X = 1, Y = 3), (X = 2, Y = 3), (X = 3, Y = 2)$$

and the car seems to be in door one half of the possibilities. However, these events do not have the same probability of occurring. However, if the host changes his strategy, the conditional probabilities fall more in line with intuition – if the host always picks door number 2 to open if door number 1 was picked and had the car behind it, then the two remaining doors have an equal chance of being picked.

We end this chapter with a final probability rule which is important in statistical analysis. If B is partitioned into a finite sequence of disjoint events A_1, \dots, A_n , then we have the formula $\mathbf{P}(B) = \sum_i \mathbf{P}(B|A_i)\mathbf{P}(A_i)$. This easily gives us Bayes rule

$$\mathbf{P}(A_j|B) = \frac{\mathbf{P}(B|A_j)}{\sum_i \mathbf{P}(B|A_i)\mathbf{P}(A_i)}$$

If we view A_j as a particular hypothesis from the set of all hypotheses, and B as some obtained data, then Bayes rule enables us to compute the probability that A_j is the true hypothesis from the probability that B is the data generated given the hypothesis is true. This is incredibly important if you can interpret these probabilities correctly (if you are a Bayesian), but not so useful if you are an empiricist (in which case we assume there is a ‘true’ result we are attempting to estimate from trials, so there is no probability distribution over the correctness hypothesis, other than perhaps a point

mass, in which case Bayes rule gives us no information). We reiterate that Bayes rule is a theorem of probability theory, so is true in any interpretation, but can be used by Bayesians in a much more applicable way to their statistical analysis.

1.5 Kolmogorov's Zero-One Law

s

1.6 Beware of Intuition

Intuition is very useful in mathematics. It hints at which paths to a solution to a problem are most likely to bear fruit. In such a practical mathematical subject in probability, this is certainly true. But probability also has many pitfalls that often trip up our minds. And it is useful to list them in one place so that we are able to better acknowledge, and avoid them.

- *The Law Of Averages*: If X is a random variable, then it is highly likely to be close to its average. For instance, a British newspaper once made the claim that to pick a 'good' set of six numbers in the lottery, the average of the numbers should be close to 25. Of course, this is false; the numbers in the lottery have no actual bearing on the likelihood of being drawn. To those who are not easily convinced, ask if the game is the same if the balls were painted with various colours, rather than inscribed with numbers. If you accept this, ask if the colors can be swapped without changing the probabilities. And so if you accept this, then the numbers on the balls can be swapped, so all draws occur with equal probability.

Of course, theoretically if $X = (X_1, \dots, X_6)$ are six independent random variables uniformly distributed in $\{0, \dots, 9\}$, and we set $S(X) = X_1 + \dots + X_6$, then

$$\mathbf{E}(S(X)) = 6 \cdot \mathbf{E}(\text{Uni}(0, \dots, 9)) = 27.$$

Certainly, $\arg\max_n \mathbf{P}(S = n) = 27$. But this does not mean that any particular choice of $(x_1, \dots, x_6) \in \{0, \dots, 9\}^6$ with $x_1 + \dots + x_6 = 27$ is

more likely to occur than (x_1, \dots, x_6) with $x_1 + \dots + x_6 \neq 27$; there are just more choices of $(x_1, \dots, x_6) \in \{0, \dots, 9\}^6$ with $x_1 + \dots + x_6 = 27$.

What further confuses matters is the weak law of large numbers, which says that if $X^1, X^2, \dots, X^n \in \{0, \dots, 9\}^6$ are n independant draws of 6 one digit numbers, for some large value n , then the random variable

$$\frac{1}{n} \sum_{i=1}^n S(X^i)$$

is *very* likely to be close to 27. But just because the average is likely to be close to 27, does not at all mean that each $S(X^i)$ need not be close to 27.

- *Assuming Independence*: This is the opposite of the law of averages, assuming certain events are not related to one another, when in fact they are. Make sure to carefully think through whether certain events are independant, before you use this in an argument.
- *Assuming Equal Likelihood*: As David Williams' says, concentration on the 'equally likely' approach to probability is an invitation to disaster.
- *The Behaviour of Ratios*: TODO.

The remainder of this section is devoted to certain famous paradoxes in probability theory.

Example. Suppose that a person has a 1% probability of contracting the disease. A test for the disease has a 90% accuracy of success. One person is chosen at random, tested for the disease, and the test comes back positive. It may seem that the person has reason to worry, but the person is actually still more likely to not have the disease, than to have the disease. If we let D the event that the person has the disease, and P the event that the test is positive, then using Bayes' rule, we conclude

$$\mathbf{P}(D|P) = \frac{\mathbf{P}(P|D) \cdot \mathbf{P}(D)}{\mathbf{P}(P|D) \cdot \mathbf{P}(D) + \mathbf{P}(P|D^c) \cdot \mathbf{P}(D^c)} = \frac{0.9 \cdot 0.01}{0.9 \cdot 0.01 + 0.1 \cdot 0.99} = 1/12.$$

So you are actually 11 times more likely to not have the disease than to have the disease.

Chapter 2

Random Variables

The formality of probability theory is ironic, because even though we require the theory of measures and real analysis to place the foundations of the theory on a rigorous footing, in the probabilistic way of thinking we try to eschew as much of this foundation as possible; we try and study properties of objects which are ‘independant’ of the sample space we consider. As a rough approximation, if $T : X \rightarrow Y$ is a surjective measure preserving map between probability spaces (X is an extension of the space Y , allowing more outcomes), then the random variable $Y \circ T$ is considered the ‘same’ as the random variable Y , and the concepts studied in probability theory should be preserved under this extension. As we reach further and further into statistical theory, samples spaces will soon become a distant memory, brought back only for the most technical of arguments. The irony of introducing the sample space is unfortunate, because while the space is in the background, in the probabilistic way of thinking about problems we try and eschew the sample space as much as possible.

Thus we emphasize the difference between how we formally define a random variable, and how we ‘think’ of a random variable. Formally, a random variable is a Borel measurable function $X : \Omega \rightarrow \mathbf{R}$. But we think of a random variable as a ‘random number’, which takes a particular value for each point in the sample space Ω . The measurability is needed so that the ‘set of questions we ask about X ’ are still legal in the probability space Ω , i.e. the quantities

$$\mathbf{P}(X \in A) = \mathbf{P}(X^{-1}(A))$$

are well defined for each Borel set A .

2.1 Distributions

A random variable X induces a probability measure on the Borel sigma algebra of \mathbf{R} , which we call the **law**, or **distribution** of X . It is defined as $\mathbf{P}_X(A) = \mathbf{P}(X^{-1}(A))$. The distribution captures the size and shape of the random variable, but not its relation to other random variables in the sample distribution.

If X is uniformly distributed on $[0, 1]$, then μ_X is the Lebesgue measure on $[0, 1]$, and if X is a *discrete random variable*, in the sense that the range of X takes on only countably many values, then μ_X is a discrete measure with $\mu_X(\{a\}) = \mathbf{P}(X = a)$.

We say two random variables X and Y are **identically distributed**, sometimes written

$$X \stackrel{(d)}{=} Y$$

if $\mu_X = \mu_Y$. Note that X and Y need not even be defined on the same sample space, let alone be actually equal to one another. For instance, if $\Omega = \{0, \dots, 6\}^2$, with the uniform measure, and X and Y are random variables with $X(a, b) = a$ and $Y(a, b) = b$, then X and Y are identically distributed, but they are not equal to one another.

2.2 Expectation

Theorem 2.1. For any $X \geq 0$,

$$\mathbf{E}[X] = \int \mathbf{P}(X \geq x) dx$$

Proof. Applying Fubini's theorem,

$$\begin{aligned} \int_0^\infty \mathbf{P}(X \geq x) dx &= \int_0^\infty \int_x^\infty d\mathbf{P}_*(y) dx \\ &= \int_0^\infty \int_0^y dx d\mathbf{P}_*(y) \\ &= \int_0^\infty y d\mathbf{P}_*(y) = \mathbf{E}[X] \end{aligned}$$

□

Chapter 3

Tail Bounds and Concentration

In most cases, it is very difficult, if not impossible, to find a closed form formula for the probability of certain events. But often, all that is important is to obtain some upper or lower bound on the probability of some *bad* event happening, or lower bound the probability of a *good* event happening. It will be helpful to utilize Tao's notation for the probabilities of certain events. Let E be an event depending on some parameter n :

- E holds *asymptotically almost surely* if $\mathbf{P}(E) = 1 - o(1)$ as $n \rightarrow \infty$.
- E holds with *high probability* if $\mathbf{P}(E) = 1 - O(n^{-c})$ for some $c > 0$.
- E holds with *overwhelming probability* if $\mathbf{P}(E) = 1 - O_c(n^{-c})$ for every $c > 0$.
- E holds *almost surely* if $\mathbf{P}(E) = 1$.
- E holds *surely* if $E^c = \emptyset$.

The union bound $\mathbf{P}(E \cup F) \leq \mathbf{P}(E) + \mathbf{P}(F)$ shows these classes are closed under various numbers of intersections:

- If $\{E_\alpha\}$ is an *arbitrary* family of events which occur surely, then $\bigcap E_\alpha$ occurs surely.
- If $\{E_k\}$ is a *countable* family of events which occur almost surely, then $\bigcap E_k$ occurs almost surely.

- If E_1, \dots, E_N are a *polynomial* number of events and these events hold with *uniform* overwhelming probability (i.e. implicit constants are independent of events), then $\bigcap E_k$ holds with overwhelming probability. In other words, if $N = O(n^{O(1)})$, and $\mathbf{P}(E_k) = 1 - O_c(n^{-c})$ independently of k , then for all c ,

$$\mathbf{P}\left(\bigcap E_k\right) = 1 - \mathbf{P}\left(\bigcup E_k^c\right) \geq 1 - N \cdot O_c(n^{-c}) = 1 - O_c(n^{O(1)-c}).$$

Since c can be arbitrarily large, this gives the result.

- If E_1, \dots, E_N are a *sub-polynomial* number of events, holding with *uniform* high probability, then $\bigcap E_k$ holds with high probability. In other words, if $N = O(n^{o(1)})$, and there is c such that $\mathbf{P}(E_k) = 1 - O(n^{-c})$ for all k , then for any $c_0 < c$,

$$\mathbf{P}\left(\bigcap E_k\right) = 1 - \mathbf{P}\left(\bigcup E_k^c\right) \geq 1 - O(n^{o(1)-c}) = 1 - O(n^{-c_0}).$$

- If E_1, \dots, E_N are a *finite* number of events, holding asymptotically almost surely, then $\bigcap E_k$ holds asymptotically almost surely. In other words, if $N = O(1)$, then

$$\mathbf{P}\left(\bigcap E_k\right) = 1 - N \cdot o(1) = 1 - o(1).$$

Notice that as an event becomes more certain, we are more free to apply larger and large intersections. Event that are likely remain likely under suitable types of condition, because

$$\mathbf{P}(F|E) = \frac{\mathbf{P}(F \cap E)}{\mathbf{P}(E)} \leq \frac{\mathbf{P}(F)}{\mathbf{P}(E)}$$

In particular, let F and E be events depending on a parameter n :

- If F occurs almost surely, it occur almost sure conditioned on E .
- If F occurs with overwhelming probability, it occurs with overwhelming probability conditioned on E provided that $\mathbf{P}(E) \gtrsim n^{-C}$ for some $C > 0$, since

$$\mathbf{P}(F^c|E) \leq \frac{\mathbf{P}(F^c)}{\mathbf{P}(E)} = \frac{O_c(n^{-c})}{O(n^{-C})} = O_c\left(n^{-(c-C)}\right)$$

- If F occurs with high probability, it occurs with high probability conditioned on E provided that $\mathbf{P}(E) \gtrsim n^{-c_0}$ for some sufficiently small c_0 depending on the decay of F . If $\mathbf{P}(F) = 1 + O(n^{-c})$, and $c_0 < c$, then

$$\mathbf{P}(F^c|E) \leq \frac{\mathbf{P}(F^c)}{\mathbf{P}(E)} \leq O\left(n^{-(c-c_0)}\right)$$

- If F occurs asymptotically almost surely, it occurs asymptotically almost surely conditioned on E , provided that $\mathbf{P}(E) \gtrsim 1$.

None of the consequence of this list involve any assumptions about the interactions of the events. Given more interesting interactions between the events, one can often strengthen the bounds obtained here.

3.1 Tail Bounds

Tail bounds upper bound the probability that a random variable will take large values. Markov's inequality is a basic version, which connects the expectation of a random variable with a tail bound.

Theorem 3.1 (Markov's Inequality). *If $X \geq 0$, then $\mathbf{P}(X \geq t) \leq \mathbf{E}(X)/t$.*

Proof. By Fubini's theorem,

$$\mathbf{E}(X) = \int \int_0^X dt d\mathbf{P} = \int_0^\infty \int \mathbf{I}(t \leq X) d\mathbf{P} dt = \int_0^\infty \mathbf{P}(X \geq t) dt$$

Markov's inequality is thus obtained as a weak-type inequality. \square

Remark. The integral identity

$$\mathbf{E}X = \int_0^\infty \mathbf{P}(X > t) dt$$

also enables one to bound $\mathbf{E}(X)$ in terms of a tail bound.

One refers to applications of Markov's inequality as a *first moment* calculation, to compare with other methods for bounding tail probabilities. In general, if we bound $\mathbf{E}(X^p)$ for some $p > 0$, then Markov's inequality shows that

$$\mathbf{P}(X \geq t) = \mathbf{P}(X^p \geq t^p) \leq \mathbf{E}(X^p)/t^p$$

As p becomes larger, this inequality gives a sharper rate of decay on the random variable. This is known as a *p th moment bound*. We can be even more clever in this respect. If $\mathbf{E}(e^{\lambda X})$ is finite for some $\lambda > 0$, then

$$\mathbf{P}(X \geq t) = \mathbf{P}(e^{\lambda X} \geq e^{\lambda t}) \leq \mathbf{E}(e^{\lambda X})e^{-\lambda t}$$

This is known as a *Chernoff bound*.

Remark. The linearity of expectation makes applying Markov's inequality quite simple. But higher order moments are non-linear, and so often one needs to analyze the distribution of X much more carefully. Chernoff bounds are even more difficult to obtain, but much more useful.

Remark. Technically, the dominated convergence theorem enables one to conclude the sharper results than given here. For instance, if $\mathbf{E}X^p < \infty$, then $\mathbf{P}(X > t) = o(1/t^p)$, and if $\mathbf{E}(e^{\lambda X}) < \infty$, then $\mathbf{P}(X > t) = o(e^{-\lambda t})$. But the dominated convergence theorem gives no control on the rates of this little o term, which makes it relatively ineffective.

We can quantify tail bounds of a random variable as follows:

- If X has bounded expectation, then for any $\varepsilon > 0$, $|X| = O(n^\varepsilon)$ with uniform high probability.
- If X has bounded p 'th moment, then $|X| = O(n^\varepsilon)$ with probability $1 - O(n^{-\varepsilon p})$.
- If $\mathbf{E}(e^{\lambda X})$ is bounded, then $|X| = O(\log^{O(1)} n)$ with overwhelming probability.

Very similar methods establish upper bounds on $\mathbf{P}(|X - \mathbf{E}X| \geq t)$, known as a *deviation inequality*. The equivalent to moment bounds in the deviation regime is Chebyshev's inequality.

Theorem 3.2 (Chebyshev's Inequality). *If X has finite variance σ^2 , then*

$$\mathbf{P}(|X - \mathbf{E}(X)| \geq t) \leq \sigma^2/t^2$$

More generally,

$$\mathbf{P}(|X - \mu| \geq t) \leq \frac{\mathbf{E}|X - \mu|^p}{t^p} \quad \text{and} \quad \mathbf{P}(|X - \mu| \geq t) \leq \mathbf{E}(e^{\lambda|X - \mu|})e^{-\lambda t}.$$

Proof. Applying Markov's inequality, we find

$$\mathbf{P}(|X - \mu| \geq t) = \mathbf{P}(|X - \mu|^2 \geq t^2) \leq \frac{\mathbf{E}|X - \mu|^2}{t^2} = \frac{\sigma^2}{t^2}.$$

The other equations was treated similarly. \square

Example. A Median for a random variable X is a quantity M such that

$$\mathbf{P}(X \geq M), \mathbf{P}(X \leq M) \leq 1/2$$

For an absolutely continuous random variable, the median always exists and is unique. If $X \geq 0$ has a median M , and finite variance σ^2 , Chebyshev's inequality implies

$$\mathbf{P}(|X - \mu| \geq \sqrt{2} \cdot \sigma) \leq 1/2$$

Thus $\mu - \sqrt{2} \cdot \sigma \leq M \leq \mu + \sqrt{2} \cdot \sigma$, so $M = \mu + O(\sigma)$. This explains why in most circumstances, the mean replaces the median in statistical estimates, because they are equivalent up to a small number of standard deviations.

When t is large, a bound on $\mathbf{P}(|X - \mathbf{E}X| \geq t)$ is known as a large deviation inequality. Chebyshev's inequality becomes useful only when t exceeds a single standard deviation, i.e. $t \geq \sigma$, and is therefore an example of a large deviation inequality. Second moment methods are often useful for bounding sums of independent random variables because if X and Y are independent, then $\mathbf{V}(X + Y) = \mathbf{V}(X) + \mathbf{V}(Y)$ is easy to calculate. We can continue establishing large deviation bounds, obtaining $\mathbf{P}(|X - \mathbf{E}X| > t) \lesssim 1/t^p$ by bounding $\mathbf{E}|X - \mathbf{E}X|^p$. We will refer to any of these inequalities as Chebyshev bounds.

3.2 Moments of Sums of Random Variables

We now study the tail inequalities for a sum $S_n = X_1 + \cdots + X_n$ of random variables, using the moment methods we have established. For simplicity, we assume all variables X_k have mean zero, so that deviation and tail inequalities agree with one another. Of course, general results then follow by translation.

Using the first moment method, which requires no assumption on the interdependence of the X_k , we find $\mathbf{E}|S_n| \leq \mathbf{E}|X_1| + \cdots + \mathbf{E}|X_n|$. Combined

with Markov's inequality, we conclude

$$\mathbf{P}(|S_n| \geq t) \leq \frac{\mathbf{E}|X_1| + \cdots + \mathbf{E}|X_n|}{t}$$

Nonetheless, this inequality can be sharp, especially given no independence between the random variables. which causes huge fluctuations in the magnitude of S_n . For instance, if all X_k are all equal to a single random variable $X \in \{-1, 1\}$, then $|S_n| = n$, so Markov's inequality is sharp when $t = n$.

Now we consider what the second moment method gives us. Note that

$$\mathbf{E} S_n^2 = \sum \mathbf{E}(X_i X_j)$$

If we assume the variables are pairwise independent, then $\mathbf{E}(X_i X_j) = 0$ for $i \neq j$, so $\mathbf{E} S_n^2 = \sum \mathbf{E}(X_i^2)$. This is just the familiar statement that variance is additive. Chebyshev's inequality then yields that

$$\mathbf{P}(|S_n| \geq t) \leq \frac{\mathbf{V}(X_1) + \cdots + \mathbf{V}(X_n)}{t^2}$$

Informally, this means

$$S_n = O(\mathbf{V}(S_n)^{1/2}) = O((\mathbf{V}(X_1) + \cdots + \mathbf{V}(X_n))^{1/2}).$$

We cannot expect to obtain a much sharper region of concentration for S_n than this, because

$$\mathbf{V}(X_1) + \cdots + \mathbf{V}(X_n) = \mathbf{V}(S_n) = \int_0^\infty 2t \mathbf{P}(|S_n| \geq t) dt.$$

Example. Assume for simplicity that $\mathbf{V}(X_k) = 1$ for all k , and $X_k = O(1)$, so $S_n = O(n)$ and $\mathbf{V}(S_n) = n$. Suppose $|S_n| < A_n$ with high probability, for some $A_n > 0$, i.e. $\mathbf{P}(|S_n| \geq A_n) \lesssim O(1/n^{1+\varepsilon})$ for some $\varepsilon > 0$. Then

$$\begin{aligned} n = \int_0^\infty 2t \mathbf{P}(|S_n| \geq t) &\lesssim \int_0^{A_n} 2t + \int_{A_n}^{O(n)} 2t \cdot n^{-c} \\ &\lesssim A_n^2 + (1/n^{1+\varepsilon})(O(n^2) - A_n^2) \\ &\lesssim (1 - 1/n^{1+\varepsilon})A_n^2 + n^{1-\varepsilon}. \end{aligned}$$

Thus

$$A_n^2 \gtrsim \frac{n - n^{1-\varepsilon}}{1 - 1/n^{1+\varepsilon}} \gtrsim n$$

so $A_n \gtrsim n^{1/2} = (\mathbf{V}(X_1) + \dots + \mathbf{V}(X_n))^{1/2}$, as we expected.

The next example shows the tail bound is sharp in particular circumstances, especially when only pairwise independence is satisfied.

Example. For each nonzero $k \in \{0, 1\}^m$, let $X_k = (-1)^{k \cdot Y}$, where $Y = (Y_1, \dots, Y_m)$ is drawn uniformly at random from the hypercube $\{0, 1\}^m$. If $k \neq 0$, then there is j such that $k_j = 1$, and then $\mathbf{E}((-1)^{k_i Y_i}) = \mathbf{E}((-1)^{Y_i}) = 0$, so

$$\mathbf{E}(X_k) = \prod_{j=1}^m \mathbf{E}((-1)^{k_j Y_j}) = 0$$

Also note $\mathbf{V}(X_k) = 1$ since $X_k^2 = 1$. If $k_i = 1$, then X_k is independent of the σ algebra generated by $\{Y_j : j \neq i\}$. This is because by looking at combinatorics on the hypercube, for any $y \in \{0, 1\}^m$, where we view the hypercube as a vector space over \mathbf{Z}_2 ,

$$\begin{aligned} \mathbf{P}(\{X_k = 1\} \cap \{Y_j = y_j : j \neq i\}) &= \mathbf{P}(\{Y_i = k \cdot y - y_i\} \cap \{Y_j = y_j : j \neq i\}) \\ &= 1/2^m = \mathbf{P}(X_k = 1) \mathbf{P}\left(\bigcap_{j \neq i} Y_j = y_j\right) \end{aligned}$$

If $l \in \{0, 1\}^m$ satisfies $l_i = 0$, then X_l is measurable with respect to $\sigma(Y_j : j \neq i)$, which implies it is independent to X_k . Thus the random variables are pairwise independent. On the other hand,

$$S = \sum_{k \in \{0, 1\}^m} X_k = 2^m \cdot \mathbf{I}(Y = 0)$$

is the sum of pairwise random variables with mean one. By pairwise independence, its variance is $2^m - 1$. Chebyshev's inequality gives

$$\frac{1}{2^m} = \mathbf{P}(S \geq 2^m) \leq \frac{1}{2^m - 1}$$

Thus the inequality is sharp here, up to constants which are insignificant for large m .

Regardless of the sharpness of the region of concentration, we can still hope to obtain sharper results for the tail bounds of random variables. To do this, we analyze higher moments of S_n . For simplicity, we now assume that each X_k has unit variance, still has mean zero, and furthermore, $|X_k| \leq C$ for some constant C . We note that if k is an even integer, then

$$\mathbf{E}|S_n|^k = \sum \{\mathbf{E}(X_{i_1} \dots X_{i_k}) : 1 \leq i_1, \dots, i_k \leq n\}$$

Suppose the random variables are k -wise independent, i.e. every collection of k random variables is independent from one another, this is fairly easy to calculate. Note that if some i_j appears only a single time in (i_1, \dots, i_k) , then $\mathbf{E}(X_{i_1} \dots X_{i_k}) = 0$. Thus at most $k/2$ distinct random variables appear in terms with nonzero expectation. If exactly $k/2$ distinct variables appear, then each random variable appears exactly twice, and since each X_k has unit variance, we conclude $\mathbf{E}(X_{i_1} \dots X_{i_k}) = 1$. More generally, if $k/2 - r$ distinct variables appear, then we use the fact that $X_k \leq K$ together with the unit variance bound to conclude $\mathbf{E}(X_{i_1} \dots X_{i_k}) \leq K^{2r}$. If N_r gives the total number of (i_1, \dots, i_k) such that $k/2 - r$ terms appear, then

$$\mathbf{E}|S_n|^k \leq \sum_{r=0}^{k/2} K^{2r} N_r$$

To estimate N_r , we note that there are $\binom{n}{k/2-r}$ ways to choose $k/2 - r$ integers from $\{1, \dots, n\}$. Once these numbers are fixed, each integer i_j must come from one of these $k/2 - r$ integers. Thus we obtain a crude bound,

$$N_r \leq \binom{n}{k/2-r} (k/2-r)^k \leq (en)^{k/2-r} (k/2)^{k/2+r},$$

where we have applied Stirling's formula. Thus

$$\mathbf{E}|S_n|^k \leq (enk/2)^{k/2} \sum_{r=0}^{k/2} (K^2 k/en)^r.$$

If $K^2 \leq n/k$, which holds for suitably large n , then we conclude

$$\mathbf{E}|S_n|^k \leq 2(enk/2)^{k/2}.$$

Thus

$$\mathbf{P}\left(|S_n| \geq t \cdot \sqrt{2nk/e}\right) \leq 2/t^k \quad \text{and} \quad \mathbf{P}\left(|S_n| \geq t\sqrt{n}\right) \leq \frac{2(enk/2)^{k/2}}{t^k}.$$

Note that as k increases, the rate of decay in t improves, but the range that S_n concentrates in grows from $O(\sqrt{n})$ to $O(\sqrt{nk})$.

Now we assume that the random variables $\{X_k\}$ are actually jointly independent. Then we can apply the last method for any k , and optimize for each value of t . In particular, if \sqrt{nk} is a small multiple of t , then we obtain

$$\mathbf{P}\left(|S_n| \geq t \cdot \sqrt{n}\right) \leq C \exp(-ct^2)$$

for some constants $c, C > 0$. Thus we obtain square-exponential decay by controlling all the individual moments of S_n . A slicker way to do this is using exponential moments.

3.3 Concentration of Measure

We have seen that given a large number of random variables X_1, \dots, X_N , with $X_k = O(1)$ almost surely for each k , then even though the random variable S_N ranges over an interval of length $O(N)$, it has a high probability of concentrating in a \sqrt{N} region. The basic intuition is that for large N , it is difficult for independent random variables to ‘work together’ to push the function too far from its mean. The tail bounds tend to be *sub-gaussian*, in the sense that S_N deviates t standard deviations from the mean with probability $O(e^{-ct^2})$. Concentration of measure is a primary tool for ensuring various events hold with overwhelming probability.

Lemma 3.3. *If X is centrally distributed and takes values in $[a, b]$, then*

$$\mathbf{E}(e^{\lambda X}) \leq e^{\lambda^2(b-a)^2/8}.$$

Proof. By scale homogeneity, we may assume $b - a = 1$, so it suffices to prove that $\mathbf{E}(e^{\lambda X}) \leq e^{\lambda^2/8}$. Write $X = \Lambda a + (1 - \Lambda)b$ for some random value $0 \leq \Lambda \leq 1$. Applying convexity,

$$e^{\lambda X} \leq \Lambda e^{\lambda a} + (1 - \Lambda)e^{\lambda b}$$

Hence

$$\mathbf{E}(e^{\lambda X}) \leq be^{\lambda a} - ae^{\lambda b} = e^{\lambda a}(b - ae^{\lambda}) = e^{L(\lambda)}$$

where $L(x) = ax + \ln(b - ae^x) = ax + \ln(a + 1 - ae^x)$. Note $L(0) = L'(0) = 0$, and if $x \geq 0$,

$$L''(x) = \frac{-a(a+1)e^x}{(1+a-ae^x)^2} \leq 1/4$$

By Taylor's theorem, we conclude $L(x) \leq x^2/8$, and so $\mathbf{E}(e^{\lambda X}) \leq e^{\lambda^2/8}$, which completes the proof of the result. \square

Theorem 3.4 (Hoeffding's Inequality). *Let X_1, \dots, X_n be centrally distributed independant random variables, with $a_i \leq X_i \leq b_i$. Then*

$$\mathbf{P}\left(\sum X_i \geq t\right) \leq e^{-2t^2/\sum(b_i-a_i)^2}$$

Proof. For any $\lambda > 0$, a Chernoff bound gives that

$$\mathbf{P}\left(\sum X_i \geq t\right) \leq e^{-\lambda t} \prod \mathbf{E}(e^{\lambda X_i})$$

The last lemma thus allows us to conclude

$$\mathbf{P}\left(\sum X_i \geq t\right) \leq e^{-\lambda t} \prod_{i=1}^n e^{\lambda^2(b_i-a_i)^2/8}$$

In particular, if

$$\lambda = \frac{t}{\sum(b_i-a_i)^2/4}$$

we obtain the required inequality. \square

Suppose that $a_i \leq X_i \leq b_i$, but X_i is actually much more likely to concentrate in a much smaller region. Then Chernoff's bound gives tighter results, provided one can give better bounds on $\mathbf{E}(e^{tX_i})$ for each i . For instance, if the X_i are $\{0,1\}$ valued Bernoulli variables with parameters p_i , and $t > \mu$, where $\mu = \sum \mathbf{E}(X_i)$. Then one calculates that

$$\mathbf{E}(e^{\lambda X_i}) = p_i e^{\lambda} + 1 - p_i = 1 + (e^{\lambda} - 1)p_i \leq \exp((e^{\lambda} - 1)p_i)$$

If we sum up over the X_i , then $\mathbf{E}(e^{\lambda S_N}) \leq \exp((e^{\lambda} - 1)\mu)$. For any $t > 0$, picking $\lambda = \ln(t/\mu)$ gives

$$\mathbf{P}(S_N \geq t) \leq e^{-\mu}(e\mu/t)^t.$$

This tends to be tighter for smaller deviation values of t when μ is small.

Example. Suppose we are performing a test of some property, which succeeds in obtaining the correct answer with probability $1/2 + \delta$, where δ is small. Then we obtain a stronger test by performing the test independantly N times, and taking the majority vote. If X_1, \dots, X_N denotes the $\{-1, 1\}$ valued outcome of the particular tests, our new test is just $\text{sgn}(S_N)$, where $S_N = X_1 + \dots + X_N$ (We also assume here that N is odd, so that S_N never equals zero). Without loss of generality, if the property we are testing is true, then $\mathbf{P}(X_n = 1) = 1/2 + \delta$, and $\mathbf{P}(X_n = -1) = 1/2 - \delta$. Thus the random variable has expectation 2δ . If $\Delta_n = X_n - 2\delta$, then Δ_n is centrally distributed, and $S_N = \sum \Delta_n + 2N\delta$. By Hoeffding's inequality, we conclude that

$$\mathbf{P}(\text{Test fails}) = \mathbf{P}(S_N \leq 0) = \mathbf{P}\left(\sum \Delta_n \leq -2N\delta\right) \leq e^{-8N\delta^2}$$

In particular, if we want a bound like $\mathbf{P}(\text{Test fails}) \leq \varepsilon$, then we need only choose $N = \Omega(\delta^{-2} \log(1/\varepsilon))$. This is why when performing a statistical test, a probability of success any amount higher than 50% is essentially comparable to a 99% probability of success.

The full assumption of indepdance is not necessary for a Chebyshev-type tail bound. In fact, it suffices that the differences $X_{i+1} - X_i$ can depend on the random variables X_1, \dots, X_i , but have mean zero once conditioned on the X_i . This means that $\{S_n\}$ is a *martingale*. Azuma's inequality gives such a bound.

Theorem 3.5. Let $\{S_k\}$ be a martingale (or super-martingale) with $|S_k - S_{k-1}| \leq a_k$ almost surely, and $S_0 = 0$. Then

$$\mathbf{P}(S_n \geq t) \leq \exp\left(\frac{-2t^2}{\sum a_i^2}\right).$$

Proof. We again consider the exponential moment method. We apply Hoeffding's lemma since $|S_n - S_{n-1}| \leq a_n$ to conclude

$$\begin{aligned} \mathbf{E}(\exp(\lambda S_n)) &= \mathbf{E}(\exp(\lambda(S_n - S_{n-1})) \exp(\lambda S_{n-1})) \\ &= \mathbf{E}(\exp(\lambda S_{n-1}) \mathbf{E}(\exp(\lambda(S_n - S_{n-1})) | S_{n-1})) \\ &\leq e^{a_n^2 \lambda^2 / 8} \mathbf{E}(\exp(\lambda S_{n-1})) \end{aligned}$$

Iterating this argument gives

$$\mathbf{E}(e^{\lambda S_n}) \leq \exp\left(\sum a_i^2 \lambda^2 / 8\right)$$

And so performing a Chernoff bound gives

$$\mathbf{P}(S_n \geq t) \leq \exp\left(\sum a_i^2 \lambda^2 / 8\right) e^{-\lambda t}$$

Picking $\lambda = 4t / \sum a_i^2$, we conclude

$$\mathbf{P}(S_n \geq t) \leq e^{-2t^2 / \sum a_i^2}. \quad \square$$

3.4 Subgaussian Random Variables

Hoeffding's inequality only applies to bounded random variables. But we should still obtain fast tail decay in other circumstances. For instance, a Gaussian distribution has tail decay comparable to Hoeffding's inequality, and since the central limit theorem gives intuition saying that scaled sums of random variables begin to behave like a normal distribution, we should expect fast tail decay for much more general families of sums. If $g \sim N(0, \sigma^2)$, we calculate

$$\mathbf{P}(g - \mu \geq y) = \int_y^\infty \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{x^2}{2\sigma^2}} dx \leq \frac{1}{y\sqrt{2\pi\sigma^2}} \int_y^\infty x e^{-\frac{x^2}{2\sigma^2}} dx = \frac{\sigma e^{-\frac{y^2}{2\sigma^2}}}{y\sqrt{2\pi}}$$

This quantity is almost always better than Chebyshev's inequality, since the ratio $1/y$, which measures the inaccuracy of our inequality, is nullified by the exponential function. We can find similar equalities for random variables which are 'dominated' by normal distributions. We call these random variables *subgaussian*.

Theorem 3.6. *The following are equivalent, for comparable constants $K_i > 0$.*

- (1) $\mathbf{P}(|X| \geq t) \leq 2\exp(-t^2/K_1^2)$.
- (2) $\|X\|_p \leq K_2\sqrt{p}$, for $p \in [1, \infty)$.
- (3) $M_{X^2}(\lambda^2) \leq \exp(K_3^2\lambda^2)$ if $|\lambda| \leq 1/K_3$.
- (4) $M_{X^2}(1/K_4^2) \leq 2$.
- (5) If, in addition, $\mathbf{E}X = 0$, $M_X(\lambda) \leq \exp(K_5^2\lambda^2)$.

If any of these equations hold, we say X is subgaussian.

Proof. We provide a proof that the inequalities implies one another, up to a constant change of coefficients.

- Suppose (1) holds. Then

$$\begin{aligned}\|X\|_p^p &= \int_0^\infty \mathbf{P}(|X| \geq \lambda^{1/p}) d\lambda \leq 2 \int_0^\infty \exp(-(\lambda^{1/p}/K_1)^2) d\lambda \\ &= 2K_1^p p \int_0^\infty \lambda^p \exp(-\lambda^2) \frac{d\lambda}{\lambda} = K_1^p p \cdot \Gamma(p/2) \leq K_1^p p (p/2)^{p/2}\end{aligned}$$

Thus $\|X\|_p \leq (p^{1/p} 2^{-1/2}) \sqrt{p} K_1 \leq (e^{1/e} 2^{-1/2}) \sqrt{p} K_1$, so we can set $K_2 \leq 2^{-1/2} e^{1/e} K_1$.

- Suppose (2) holds. Using Stirling's approximation, we compute

$$\begin{aligned}M_{X^2}(\lambda^2) &= \sum_{k=0}^\infty \frac{\mathbf{E}(X^{2k}) \lambda^{2k}}{k!} = \sum_{k=0}^\infty \frac{\|X\|_{2k}^{2k} \lambda^{2k}}{k!} \leq \sum_{k=0}^\infty k^k \frac{(2K_2^2 \lambda^2)^k}{k!} \\ &\leq \sum_{k=0}^\infty k^k \frac{(2K_2^2 \lambda^2)^k}{(2\pi)^{1/2} k^{k+1/2} e^{-k}} \leq \sum_{k=0}^\infty (2eK_2^2 \lambda^2)^k\end{aligned}$$

For $|\lambda| \leq 1/2e^{1/2}K_2$, since $(1-x)^{-1} \leq \exp(2x)$ if $0 \leq x \leq 1/2$, we conclude

$$M_{X^2}(\lambda^2) = \frac{1}{1 - 2eK_2^2 \lambda^2} \leq \exp(4eK_2^2 \lambda^2)$$

Thus we can set $K_3 \leq 2e^{1/2}K_2$.

- The fact that (3) implies (4) is very simple, since $M_{X^2}(\log(2)/K_3^2) \leq \exp(\log(2)) = 2$. Thus we can set $K_4 \leq K_3/\log(2)^{1/2}$.
- The fact that (4) implies (1) is also pretty simple, since a Chernoff type bound gives

$$\mathbf{P}(|X| \geq t) = \mathbf{P}(e^{|X|^2/K_4^2} \geq e^{t^2/K_4^2}) \leq M_{X^2}(1/K_4^2) e^{-t^2/K_4^2} \leq 2e^{-t^2/K_4^2}$$

So we can set $K_1 \leq K_4$. This completes the argument that (1) through (4) are all equivalent.

- Now we show (3) implies (5), assuming X is centrally distributed. If $|\lambda| \leq 1/K_3$, we use the inequality $e^{\lambda x} \leq \lambda x + e^{\lambda^2 x^2}$ to conclude

$$M_X(\lambda) \leq \lambda \mathbf{E}X + M_{X^2}(\lambda^2) = M_{X^2}(\lambda^2) \leq \exp(K_3^2 \lambda^2)$$

If $|\lambda| \geq 1/K_3$, we use the inequality $\lambda x \leq \lambda^2/2 + x^2/2$, so that

$$M_X(\lambda) \leq e^{K_3^2 \lambda^2/2} M_{X^2}(1/2K_3^2) \leq e^{K_3^2 \lambda^2/2} e^{1/2} \leq e^{K_3^2 \lambda^2}$$

Thus we may set $K_5 \leq K_3$.

- Suppose (5) holds. Then a Chernoff bound gives

$$\mathbf{P}(X \geq t) \leq M_X(\lambda) e^{-\lambda t} \leq e^{K_5^2 \lambda^2 - \lambda t}$$

If we set $\lambda = t/2K_5^2$, we find $\mathbf{P}(X \geq t) \leq e^{-t^2/4K_5^2}$. Thus $\mathbf{P}(|X| \geq t) \leq 2e^{-t^2/4K_5^2}$, and so we can set $K_1 \leq 2K_5$.

In particular, if we check all the constants, we see that $K_i \leq 10K_j$ for any minimal choice of K_i and K_j . \square

The natural way to form a measure of being subgaussian is to come up with an Orlicz norm. Given a convex, increasing function $\psi : [0, \infty) \rightarrow [0, \infty)$, with $\psi(0) = 0$, and $\psi(x) \rightarrow \infty$ as $x \rightarrow \infty$. We define the Orlicz norm $\|X\|_\psi$ corresponding to ψ as the infimum of t such that $\mathbf{E}(\psi(t^{-1}|X|)) \leq 1$.

Theorem 3.7. *The Orlicz norm is actually a norm, and is also complete.*

Proof. It is obvious that $\|\lambda X\|_\psi = |\lambda| \|X\|_\psi$. To obtain the triangle inequality, we note that if $\mathbf{E}(\psi(t^{-1}|X|)) \leq 1$, and $\mathbf{E}(\psi(s^{-1}|Y|)) \leq 1$, then convexity implies

$$\mathbf{E} \left(\psi \left(\frac{|X+Y|}{s+t} \right) \right) \leq \mathbf{E} \left(\frac{|X|+|Y|}{s+t} \right) \leq \frac{t \mathbf{E}(t^{-1}|X|) + s \mathbf{E}(s^{-1}|Y|)}{s+t} \leq 1$$

This gives $\|X+Y\|_\psi \leq 1$. If $\|X\|_\psi = 0$, then the increasing nature of ψ makes it easy to see that $X = 0$ almost surely. If X_1, X_2, \dots is a Cauchy sequence with respect to the Orlicz norm, we may thin the sequence out so that $\|X_{n+1} - X_n\|_\psi \leq 1/2^n$. This means that $S = \sum |X_{n+1} - X_n|$ has finite Orlicz norm, because if $S_N = \sum_{n \leq N} |X_{n+1} - X_n|$, then $\|S_N\|_\psi \leq 1$, and by

the monotone convergence theorem $\mathbf{E}(\psi(S)) = \lim \mathbf{E}(\psi(S_N)) \leq 1$. Thus $\|S\|_\psi \leq 1$. This means that S is finite almost everywhere, so $\sum X_{n+1} - X_n$ converges absolutely almost everywhere. Thus X_n converges almost everywhere to some X , and monotone convergence implies X_n converges to X in the Orlicz norm. But because the original sequence is Cauchy, this means the original sequence also converges to X . \square

We can use the Orlicz norms to provide a norm measuring how subgaussian a random variable is. It is obtained by considering the convex, increasing function $\psi_2(x) = e^{x^2}/2$. Naturally, we let $\|X\|_{\psi_2}$ denote this norm. Because $\|X\|_{\psi_2} < \infty$ if and only if X is subgaussian, we think of the ψ_2 norm as being a measure of how ‘subgaussian’ a random variable is.

For the hyperplane of centered subgaussian random variables, we can use an alternate norm. For a subgaussian random variable X with $\mathbf{E}(X) = 0$, we let $\sigma(X)$ denote the smallest value σ such that $M_X(\lambda) \leq \exp((\sigma\lambda)^2/2)$. Of course, $\sigma(X)$ and $\|X\|_{\psi_2}$ are comparable to one another, in the sense that

$$\sigma(X) \leq \|X\|_{\psi_2} \leq 10\sigma(X)$$

But utilizing σ often makes certain calculations simpler. For instance, if X is centered, it’s moment generating function is globally defined, so we can write

$$M_X(\lambda) = \mathbf{E}(\exp(\lambda X)) = \sum_{k=0}^{\infty} \mathbf{E}(X^k) \frac{\lambda^k}{k!} = 1 + \mathbf{V}(X)\lambda^2/2 + O(\lambda^3).$$

On the other hand,

$$\exp(\sigma(X)^2\lambda^2) = \sum_{k=0}^{\infty} \sigma(X)^{2k} \frac{\lambda^{2k}}{k!} = 1 + \sigma(X)^2\lambda^2 + O(\lambda^4).$$

Since $M_X(\lambda) \leq \exp(\sigma(X)^2\lambda^2)$ for all λ , if we take $\lambda \rightarrow 0$, we conclude that $\mathbf{V}(X) \leq \sigma(X)^2$.

It is obvious that σ is a homogenous quantity. To prove it is a norm, it therefore suffices to prove a triangle inequality. If $\sigma(X+Y) \leq \sigma(X) + \sigma(Y)$, we apply Hölder’s inequality to conclude that if $p^{-1} + q^{-1} = 1$,

$$\begin{aligned} M_{X+Y}(\lambda) &= \mathbf{E}(e^{\lambda X} e^{\lambda Y}) \leq \mathbf{E}(e^{p\lambda X})^{p^{-1}} \mathbf{E}(e^{q\lambda Y})^{q^{-1}} \\ &\leq \exp(p^{-1}(p\lambda\sigma(X))^2 + q^{-1}(q\lambda\sigma(Y))^2/2) \\ &\leq \exp((p\sigma(X)^2 + q\sigma(Y)^2)\lambda^2/2) \end{aligned}$$

Choosing $p = 1 + \sigma(X)/\sigma(Y)$ gives $\sigma(X + Y) \leq \sigma(X) + \sigma(Y)$. If X and Y are independent, we actually conclude that $X + Y$ is $\sqrt{\sigma^2 + \tau^2}$ subgaussian, because we needn't apply Hölder's inequality, instead computing

$$M_{X+Y}(\lambda) = M_X(\lambda)M_Y(\lambda) \leq \exp(\lambda^2(\sigma^2 + \tau^2)/2)$$

Thus we get a better tail bound on the sum.

Example. Let X_1, \dots, X_n be independent subgaussian random variables. Let $S = \sum X_k/N$. Then $\mathbf{E}(S) = \sum \mathbf{E}(X_k)/N$, and $\sigma(S) \leq \sqrt{\sum \sigma(X_k)^2}/N$. We therefore conclude that

$$\mathbf{P}(|S - \mathbf{E}(S)| \geq t) \leq 2 \exp\left(\frac{-N^2 t^2}{\sum \sigma(X_k)^2}\right)$$

Provided that $\sigma(X_k) = O(1)$, we conclude that $\mathbf{P}(|S - \mathbf{E}S| \geq t) \leq 2 \exp(-cNt)$ for some constant c , which is a very useful deviation inequality for sums of random variables.

If $\sigma(X) = 0$, then $M_X(\lambda) \leq 1$ for all λ , which gives $\mathbf{P}(|X| \geq t) \leq e^{-\lambda t}$ for all $\lambda > 0$, and taking $\lambda \rightarrow \infty$ gives $\mathbf{P}(|X| \geq t) = 0$ for all $t > 0$. Thus $X = 0$ almost surely. If X is not centered, we still let $\sigma(X)$ denote $\sigma(X - \mathbf{E}X)$. Then σ is no longer a norm on the space of subgaussian random variables, but a seminorm, whose kernel is the space of random variables constant almost surely. The next lemma shows that $\sigma(X) \lesssim \|X\|_{\psi_2}$, but if the mean of X is unbounded, it is no longer true that $\|X\|_{\psi_2} \lesssim \sigma(X)$.

Lemma 3.8 (Centering). If X is subgaussian, $\|X - \mathbf{E}X\|_{\psi_2} \lesssim \|X\|_{\psi_2}$.

Proof. We use the triangle inequality to write

$$\|X - \mathbf{E}X\|_{\psi_2} = \|X\|_{\psi_2} + \|\mathbf{E}X\|_{\psi_2} \lesssim \|X\|_{\psi_2} + |\mathbf{E}X|$$

Thus it suffices to prove $|\mathbf{E}X| \lesssim \|X\|_{\psi_2}$. But by Jensen's inequality,

$$|\mathbf{E}X| \leq \mathbf{E}|X| = \|X\|_1 \leq \|X\|_{\psi_2}$$

This completes the argument. \square

Theorem 3.9. If X_1, \dots, X_m are subgaussian random variables with $\|X_i\|_{\psi_2} \leq K$ for all i , then

$$\mathbf{E}(\max(|X_1|, \dots, |X_m|)) \lesssim K(\log m)^{1/2}.$$

Proof. Without loss of generality, assume $X_i \geq 0$ for all i . A union bound then gives

$$\mathbf{P}(\max(X_1, \dots, X_m) \geq t) \leq \sum_{i=1}^m \mathbf{P}(X_i \geq t) \leq m \cdot \exp(-t^2/K^2).$$

Thus

$$\begin{aligned} \mathbf{E}(\max(X_1, \dots, X_m)) &= \int_0^\infty \mathbf{P}(\max(X_1, \dots, X_m) \geq t) dt \\ &= K \log m + m \int_{K \log m}^\infty \exp(-t^2/K^2) \\ &\lesssim K \log m + Km \log m \exp(-c(\log m)^2) \lesssim K \log m. \quad \square \end{aligned}$$

Example. If X is a symmetric Bernoulli random variable with

$$\mathbf{P}(X = 1) = \mathbf{P}(X = -1) = 1/2$$

Then $\mathbf{E}(e^{X^2/t^2}) = e^{1/t^2}$. This shows that $\|X\|_{\psi_2} = (\log 2)^{-1/2}$

Example. If X is uniformly distributed on $[-1, 1]$, then

$$\mathbf{E}[X^k] = \frac{1}{2} \int_{-1}^1 x^k dx = \frac{1 - (-1)^{k+1}}{2(k+1)} = \begin{cases} \frac{1}{k+1} & k \text{ even} \\ 0 & k \text{ odd} \end{cases}$$

So

$$\mathbf{E}[e^{\lambda X}] = \sum_{k=0}^\infty \frac{\lambda^{2k}}{(2k+1)(2k)!} \leq \sum_{k=0}^\infty \frac{\lambda^{2k}}{2^k k!} = e^{\lambda^2/2}$$

so $\sigma(X) \leq 1$. Since $\mathbf{V}(X) = 1$, we must have $\sigma(X) \geq 1$, so we actually have $\sigma(X) = 1$. More generally, by scaling, if X is uniformly distributed on $[-N, N]$, then $\sigma(X) = N$.

Example. Suppose a centrally distributed random variable X satisfies $|X| \leq M$ almost surely, then $\sigma(X) \leq M$. Assume without loss of generality that $M = 1$. Set $Y = X + 1$, and $f(t) = (e^{2t} + 1)/2 - \mathbf{E}(e^{tY})$. Since $\mathbf{E}(Y) = 1$, $f'(t) = \mathbf{E}(Y[e^{2t} - e^{tY}])$. Since $Y \leq 2$ almost surely, $f'(t) \geq 0$, and so f is increasing. In particular, $f(0) = 1 - 1 = 0$, so that for $t \geq 0$,

$$\mathbf{E}(e^{tX}) = e^{-t} \mathbf{E}(e^{tY}) \leq \frac{e^t + e^{-t}}{2} \leq e^{t^2/2}$$

Since we can perform the same argument for $-X$, X is 1 subgaussian.

If X is subgaussian, it ‘looks’ like a normal distribution if $\mathbf{V}(X)^2$ is approximately equal to $\|X\|_{\psi_2}^2$. This fails for certain discrete distributions. For instance, a Bernoulli random variable $X \in \{0, 1\}$ with $\mathbf{P}(X = 1) = p$ has $\mathbf{V}(X) = \mathbf{E}(X^2) - \mathbf{E}(X)^2 = p(1 - p)$, whereas $M_X(\lambda) = 1 + pe^\lambda$, which shows $\|X\|_{\psi_2} = 1/\log(1/p)$, which decreases at an arbitrarily fast rate as $p \rightarrow 0$. Thus for small values of p , $\|X\|_{\psi_2}$ is much larger than $\mathbf{V}(X)$. On the other hand, for sums of independent subgaussian random variables, these two quantities do agree with one another. This is Khintchine’s inequality.

Theorem 3.10 (Khintchine). *Let X_1, \dots, X_n be independent, mean zero subgaussian random variables. Let $S = \sum X_k$. Then for any $p \in [2, \infty)$,*

$$\mathbf{V}(S)^{1/2} \leq \|S\|_p \lesssim K p^{1/2} \mathbf{V}(S)^{1/2},$$

where $K = \max \|X_i\|_{\psi_2}$.

Proof. TODO □

Thus if $K = O(1)$, this means $\|S\|_{\psi_2}^2$ is comparable to $\mathbf{V}(S)$, so in a subgaussian sense S does ‘look’ like a normal distribution.

3.5 Subexponential Random Variables

The class of subgaussian random variables is very flexible, but some distributions just don’t have that thin a tail. There are obviously random variables like the Cauchy distribution, whose averages do not settle down asymptotically whatsoever, but there are still distributions which do seem rather well behaved, possessing moments of all orders, while not lying in the category of subgaussian random variables. In this section we consider a slightly more general category of random variables for which we can get a tail bound, the subexponential random variables.

Theorem 3.11. *The following properties are equivalent, up to changes of constants:*

- $\mathbf{P}(|X| \geq t) \leq 2 \exp(-t/K_1)$.
- $\|X\|_p \leq K_2 \cdot p$.
- $\mathbf{E}(\exp(\lambda|X|)) \leq \exp(K_3 \cdot \lambda)$ for $0 \leq \lambda \leq 1/K_3$.

- For some K_4 , $\mathbf{E}(\exp(|X|/K_4)) \leq 2$.
- If $\mathbf{E}X = 0$, $M_X(\lambda) \leq \exp(K_5^2 \lambda^2)$ for $|\lambda| \leq 1/K_5$.

If the properties hold, we say that X is **subexponential**.

We leave the equivalence to the reader. A natural norm to place on this space is the Orlicz norm induced by $\psi_1(x) = e^{x/t}/2$.

Lemma 3.12. *A variable X is subgaussian if and only if X^2 is subexponential. The product of two subgaussian random variables is subexponential. More precisely,*

$$\|XY\|_{\psi_1} \leq \|X\|_{\psi_2} \|Y\|_{\psi_2}.$$

The analogy of Hoeffding's inequality for subgaussian functions is Bernstein's inequality, which describes the tail behaviour of a sum of subexponential random variables.

Theorem 3.13. *If X_1, \dots, X_N are independant, centrally distributed, subexponential random variables. Then*

$$\mathbf{P}\left(\left|\frac{1}{N} \sum_{n=1}^N X_n\right| \geq t\right) \leq 2 \exp\left(-c \min\left(\frac{N^2 t^2}{\sum \|X_m\|_{\psi_1}^2}, \frac{Nt}{\max \|X_m\|_{\psi_1}}\right)\right)$$

An important way to think of Bernstein's inequality is that it places the behaviour of deviation from the mean into two categories. For small deviations from the mean, the tail bound looks like that of a normal distribution. On the other hand, for large deviations, the tail becomes heavier, like an exponential distribution. The reason for the appearance of the exponential decay is a heuristic result of the central limit, which says the sum should look approximately like a normal distribution for large N .

Chapter 4

Existence Theorems

In certain fields of probability theory, we wish to discuss collections of random variables defined over the same sample space. For instance, given a sequence $\mathbf{P}_1, \mathbf{P}_2, \dots, \mathbf{P}_n$ of probability distributions defined over a space Y , we may want to talk about a sequence of independent random variables $X_i : \Omega \rightarrow Y$, such that $\mathbf{P}(X_i \in U) = \mathbf{P}_i(U)$. The construction here is simple; we take $\Omega = Y^n$, let $X_i = \pi_i$ be the projection on the i 'th variable, and let \mathbf{P} be the probability measure induced by

$$\mathbf{P}(U_1 \times U_2 \cdots \times U_n) = \mathbf{P}_1(U_1)\mathbf{P}_2(U_2)\dots\mathbf{P}_n(U_n)$$

The construction here is simple because we have finitely many distributions, but the problem becomes much harder when we need to talk about an infinite family of distributions \mathbf{P}_i , or when we need to talk about non-independent random variables, with some specified relationships between the variables. The problem is to show there exists a sample space Ω 'big enough' for the random variables to all be defined on the space.

Chapter 5

Entropy

Let μ be a probability distribution. We would like to measure the expected ‘amount of information’ contained in the distribution – in essence, the average information entropy of μ . It was Claude Shannon who found the correct formula to measure this.

Shannon considered the problem of efficient information transfer. Suppose there was a channel of communication between two friends A and B . The friends have agreed on a standard dictionary X of possible messages, along with a probability distribution μ over the dictionary, and we would like to encode these messages into bits, in such a way that the average length of the message is smallest. We then define this to be the information entropy of μ . Shannon showed that if μ is discrete with probabilities p_1, \dots, p_n , then the entropy can be calculated as

$$H(\mu) = \sum p_n \log_2 \left(\frac{1}{p_n} \right)$$

where the entropy is measured in bits, we can define the entropy in terms of the natural logarithm, in which case the entropy is said to be measured in nats. We assume that $p_i \log 1/p_i = 0$ for $p_i = 0$, which makes sense by the continuity of $x \log(1/x)$.

The entropy of a distribution also tells us

Now suppose that we were attempting to optimize a message with respect to a discrete distribution μ , and we instead encounter a distribution ν . Then the policy we have used for messages will be less optimal than if we had known that ν was the distribution in the first place. We define the relative difference in information between μ and ν as the difference

between the encoding of ν with respect to μ , and the encoding of μ with respect to μ . This is not a linearly ordered relation, ν does not possess more information than μ , just different information. If μ takes probabilities p_i and ν takes relative probabilities q_i , the difference in information is calculated to be

$$D(\mu, \nu) = \sum p_i \log(1/q_i) - \sum p_i \log(1/p_i) = \sum p_i \log(p_i/q_i)$$

This is known as the **Kullback Leibler distance** between μ and ν .

Now suppose we are viewing independent samples X_1, \dots, X_n , but we do not know where the samples are drawn from μ or ν . The larger $D(\mu, \nu)$ is, the less time we should take to make an accurate decision that the distribution is μ or ν . Indeed, if $p_i > 0$ and $q_i = 0$, then $D(\mu, \nu) = \infty$, and we can conclude with certainty that the distribution is μ if we ever view the outcome corresponding to p_i .

It is necessary to define the ‘entropy’ of an arbitrary distribution, but it is then not clear how to interpret the entropy, since an encoding of uncountably many values will always have an infinite expected number of bits. However, we can define the relative entropy by performing a discretization; Let μ and ν be distributions on some sample space X . Consider function $f : X \rightarrow \{1, \dots, n\}$, and define

$$D(\mu, \nu) = \sup_f D(f_*\mu, f_*\nu)$$

where f_* pushes measures on X onto discrete measures on $\{1, \dots, n\}$. For a fixed f , $D(\mu, \nu)$ upper bounds the difference in information we expect to see over a particular discretization. One can then calculate that

$$D(\mu, \nu) = \begin{cases} \infty & : \mu \not\ll \nu \\ \int \log(d\mu/d\nu) d\mu & : \mu \ll \nu \end{cases}$$

The relative entropies of well known distributions are easy to compute. Normal distributions, for instance, have

$$D(N(\mu_1, \sigma^2), N(\mu_2, \sigma^2)) = (\mu_1 - \mu_2)^2 / 2\sigma^2$$

For Bernoulli distributions, we have

$$D(B(p), B(q)) = p \log(p/q) + (1-p) \log\left(\frac{1-p}{1-q}\right)$$

Which is true except perhaps at boundary conditions.

The Kullback Leibler distance gives us certain bounds which are essential to information theoretic lower bounds. The bound is useful, for it relates the probabilities of distributions by the difference in information contained within.

Theorem 5.1 (The High Probability Pinsker Bound). *If μ and ν are probability measures on the same space X , and $U \subset X$ is measurable, then*

$$\mu(A) + \nu(A^c) \geq \frac{1}{2} e^{-D(\mu, \nu)}$$

Suppose we have a decision procedure which attempts to distinguish between events in probability distributions. If we choose an event A upon which the decision procedure fails to make the correct decision on the measure μ , and A^c measures the decision to fail under the measure ν , then the bound above shows the decision procedure cannot work reliably on both μ and ν .

Chapter 6

Appendix: Uniform Integrability

Uniform integrability provides stronger conditions on controlling convergence in the L^1 norm. For $p > 1$, inequalities often have ‘smoothing’ properties that are not apparent for the $p = 1$ case, so uniform integrability provides particular techniques to help us. We start with a basic result in measure theory, specialized to probabilistic language.

Lemma 6.1. *If $X \in L^1(\Omega)$ is a random variable, then for any $\varepsilon > 0$, there is $\delta > 0$ such that for any event E with $\mathbf{P}(E) \leq \delta$,*

$$\int_E |X| < \varepsilon$$

Proof. Suppose that there exists some ε , and events E_1, E_2, \dots with $\mathbf{P}(E_k) \leq 1/2^k$ but with

$$\int_{E_k} |X| \geq \varepsilon$$

By taking successive unions, we may assume the E_i are a decreasing family of sets, and then

$$\int_{\bigcap_{k=1}^{\infty} E_k} |X| = \lim_{k \rightarrow \infty} \int_{E_k} |X| \geq \varepsilon$$

and $\mathbf{P}(\bigcap E_k) = 0$, which is impossible. □

Corollary 6.2. *If $X \in L^1(\Omega)$, and $\varepsilon > 0$, then there is $K \in [0, \infty)$ with*

$$\int_{|X| > K} |X| < \varepsilon$$

A family of random variables $\{X_\alpha\}$ is called **uniformly integrable** if given $\varepsilon > 0$, there is $K \in [0, \infty)$ such that

$$\int_{|X_\alpha| > K} |X_\alpha| < \varepsilon$$

so that we can uniformly control the integral of X_α over large sets. We note that

$$\mathbf{E}|X_\alpha| = \int_{|X_\alpha| > K} |X_\alpha| + \int_{|X_\alpha| \leq K} |X_\alpha| \leq \varepsilon + K$$

so a family of uniformly integrable random variables is automatically in $L^1(\Omega)$, and *bounded* in $L^1(\Omega)$. However, a family of random variables bounded in $L^1(\Omega)$ is *not* necessarily uniformly integrable.

Example. Let Ω be $[0, 1]$ together with the Lebesgue measure. Let $E_n = (0, 1/n)$, and set $X_n = n\chi_{E_n}$. Then the X_n are bounded in $L^1(\Omega)$, but for $n \geq K$,

$$\int_{X_n > K} X_n = 1$$

and so the random variables are not uniformly integrable.

These ‘concentrating bumps’ are essentially the only reason why we cannot always exchange expectations and limits, and require the application of the dominated convergence theorem. The condition of uniform integrability removes the ability for concentrating bumps to hide within the expectation of a family of random variables, and we find it also gives us conditions that guarantee we can exchange limits with integration. First, note that if we take a concentrated bump function $X = n\chi_{E_n}$, then $\mathbf{E}|X| = 1$ is bounded uniformly over n , but $\mathbf{E}|X|^{1+\varepsilon} = n^\varepsilon$ is unbounded, reflecting the fact that boundedness in $L^p(\Omega)$ for $p > 1$ removes concentrated bump functions by magnifying their effect.

Theorem 6.3. Suppose that $\{X_\alpha\}$ is a class of random variables bounded in L^p for $p > 1$, then $\{X_\alpha\}$ is uniformly integrable.

Proof. Consider some $A \in [0, \infty)$ which gives a uniform bound $\mathbf{E}|X_\alpha|^p < A$. Applying Hölder’s inequality, we conclude

$$\int_{|X_\alpha| > K} |X_\alpha| \leq \int_{|X_\alpha| > K} \frac{|X_\alpha|^p}{K^{p-1}} \leq \frac{A}{K^{p-1}}$$

This is a uniform bound, and we may let $K \rightarrow \infty$ to let the bound go to zero. Thus the family $\{X_\alpha\}$ is uniformly integrable. \square

Corollary 6.4. *If $|X_\alpha| \leq Y$ is a uniform bound over a family $\{X_\alpha\}$ of random variables, where $Y \in L^1(\Omega)$, then $\{X_\alpha\}$ is uniformly integrable.*

Proof. We find

$$\int_{|X_\alpha| > K} |X_\alpha| \leq \int_{|X_\alpha| > K} Y \leq \int_{Y > K} Y$$

and as $K \rightarrow \infty$, $\mathbf{P}(Y > K) \rightarrow 0$, and we can apply the continuity result to conclude that

$$\int_{Y > K} Y \rightarrow 0$$

and so we obtain a uniform bound. \square

We recall that a sequence X_1, X_2, \dots of random variables **converges in probability** to a random variable X if, for every ε , $\mathbf{P}(|X_n - X| > \varepsilon) \rightarrow 0$. If $X_i \rightarrow X$ almost surely, then $X_i \rightarrow X$ in probability, because we can apply the reverse Fatou lemma to conclude

$$\begin{aligned} 0 &= \mathbf{P}\left(\liminf_{n \rightarrow \infty} |X_n - X| > \varepsilon\right) \\ &= \mathbf{P}(\limsup\{|X_n - X| > \varepsilon\}) \geq \limsup \mathbf{P}(|X_n - X| > \varepsilon) \end{aligned}$$

Hence $\mathbf{P}(|X_n - X| > \varepsilon) \rightarrow 0$. The bounded convergence theorem links L^1 convergence to convergence in probability using uniform integrability.

Theorem 6.5. *If X_n is a sequence of bounded random variables which tend to a random variable X in probability, then $X_n \rightarrow X$ in the L^1 norm.*

Proof. Let us begin by proving that if $|X_n| \leq K$, then $|X| \leq K$ almost surely. This follows because for any k ,

$$\mathbf{P}(|X| > K + 1/k) \leq \mathbf{P}(|X - X_n| > 1/k) \rightarrow 0$$

so $\mathbf{P}(|X| > K + 1/k) = 0$, and letting $k \rightarrow \infty$ gives $\mathbf{P}(|X| > K) = 0$. Let $\varepsilon > 0$ be given. Then if we choose n large enough that

$$\mathbf{P}(|X_n - X| > \varepsilon) \leq \varepsilon$$

then

$$\begin{aligned}\mathbf{E}|X_n - X| &= \int_{|X_n - X| > \varepsilon} |X_n - X| + \int_{|X_n - X| \leq \varepsilon} |X_n - X| \\ &\leq 2K\varepsilon + \varepsilon\end{aligned}$$

we can then let $\varepsilon \rightarrow 0$ to obtain L^1 convergence. \square

All this discussion concludes with a sufficient condition for L^1 convergence, showing that uniform integrability is really the right condition which removes the pathologies which prevent us from exchanging expectation with pointwise limits.

Theorem 6.6. *Let X_n be a sequence of integrable random variables, and X is another integrable random variable. Then $X_n \rightarrow X$ in the L^1 norm if and only if $X_n \rightarrow X$ in probability, and $\{X_n\}$ is uniformly integrable.*

Proof. Fix $K > 0$, and consider

$$f_K(x) = \begin{cases} K & : x > K \\ x & : |x| \leq K \\ -K & : x < -K \end{cases}$$

Then for every $\varepsilon > 0$, we can choose K such that $\|f_K(X_n) - X_n\|_1 \leq \varepsilon$, $\|f_K(X) - X\|_1 \leq \varepsilon$ uniformly across n (adding a single variable to a uniformly integrable random variable keeps it uniformly integrable). But it is easy to see that $f_K(X_n) \rightarrow f_K(X)$ in probability also, so by the bounded dominated convergence theorem, we conclude that $\|f_K(X_n) - f_K(X)\| \rightarrow 0$. A triangle inequality result gives the general result because the behaviour of X for large values is bounded by the uniform integrability.

To verify the reverse condition, note that if $\mathbf{E}|X_n - X| \rightarrow 0$, then Markov's inequality gives

$$\mathbf{P}(|X_n - X| \geq K) \leq \frac{\mathbf{E}|X_n - X|}{K} \rightarrow 0$$

to obtain uniform integrability, note that for each n , $\{X_1, \dots, X_n, X\}$ is uniformly integrable, then for each $\varepsilon > 0$, there is δ such that if $\mathbf{P}(E < \delta)$,

$$\int_E |X_n| < \varepsilon \quad \int_E |X| < \varepsilon$$

Since the entire set of X_n are bounded in $L^1(\Omega)$, we can choose K such that $\sup \mathbf{E}|X_k| < \delta K$, and then for $m > n$, $\mathbf{P}(|X_m - X| > K) < \delta$, and so

$$\int_{|X_m| > K} |X_m| \leq \int_{|X_m| > K} |X| + \mathbf{E}|X - X_m| \leq 2\varepsilon$$

where we assume we have chosen n large enough that $\mathbf{E}|X - X_m| \leq \varepsilon$. The fact that for $m \leq n$,

$$\int_{|X_m| > K} |X_m| \leq \varepsilon$$

follows from uniform integrability of the family $\{X, X_1, \dots, X_n\}$, so we have shown the entire infinite sequence is uniformly integrable. \square

Chapter 7

Percolation Theory

Let us consider the two dimensional theory of percolation. The two examples we have in mind are the lattice \mathbf{Z}^2 , and the triangular lattice \mathbf{T} . For any $p \in [0, 1]$, we define a graph structure on \mathbf{Z}^2 , adding an edge between two adjacent elements of the lattice with independent probability p . On \mathbf{T} , we instead consider *site percolation*, where we keep a hexagon with probability p .

Theorem 7.1 (Russo-Seymour-Welsh). *If $p = 1/2$, then for any $a, b > 0$, there exists c such that if A_n denotes the event that we can travel from the left edge to the right edge of the lattice $[0, a \cdot n] \times [0, b \cdot n] \cap \mathbf{Z}^2$, then $c < \mathbf{P}(A_n) < 1 - c$.*

One of the main problems in percolation theory is to determine how likely it is to find an infinite connected set of vertices, or cluster, in the randomly selected graph. As the probability of each edge becomes more likely, the graph becomes more and more connected. We find that for $p > 1/2$, there is almost surely an infinite cluster, and for $p < 1/2$, there is almost surely *not* a cluster. The value $p = 1/2$ is therefore called the *phase transition* point. A very related value to the phase transition problem is the percolation density function θ , which for each p , gives the probability $\theta(p)$ of the origin being in an infinite cluster of the graph. As an example, it is known that on \mathbf{T} , $\theta(p) = (p - 1/2)^{5/36 + o(1)}$, as $p \downarrow 1/2$. Determining phase transition points is the main focus of this chapter's notes.

7.1 Duality

Note an important duality in these geometric scenarios. Given any graph on \mathbf{Z}^2 , we can obtain another graph, the dual graph, by taking the vertices as unit squares with corners on \mathbf{Z}^2 , and with an edge between adjacent squares if there is no edge separating the two squares. Then the probability that there is an edge between two squares is the same as the probability that we do *not* select the corresponding separating edge, i.e. with probability $1 - p$. This will be useful.

The book says the dual graph of T is T , but I don't quite understand why?

7.2 Boolean Functions and Sharp Thresholds

If G is a graph, then the family of all graph structures on these vertices can be identified with $\{0,1\}^E = \{0,1\}^{O(V^2)}$. Thus a function f on the set of graphs can be identified with a boolean function, and we can apply boolean function techniques. In our case, the natural graphs will be the subgraphs of \mathbf{Z}^2 of the form $[0, a \cdot n] \times [0, b \cdot n]$. An example of a Boolean function on graphs is obtained by setting

$$f_n(G) = \mathbf{I}(\text{There is a path from left to right in } G)$$

Boolean analysis tells us the main features of G . Here we introduce some basics, which will help us get the job done.

If f is a boolean function, we say it is monotone if $x_i \leq y_i$ for each i implies $f(x) \leq f(y)$. An index i is pivotal for an input x if $f(x) \neq f(x^i)$, where x^i is x with the bit flipped at the i 'th position. The influence $\mathbf{I}_i(f)$ of f in the variable i is then the probability that for a randomly chosen x , i is pivotal. If we instead choose an input to be equal to one with probability p , and zero with probability $1 - p$, then the probability that i is pivotal is denoted $\mathbf{I}_i^p(f)$. The sum of the influence over all influences i is known as the *total influence*. Now if E is a monotone event, then it is obvious that as p increases, $\mathbf{P}(E)$ should increase. The degree to which it increases is quantified by the Margulis-Russo formula.

Theorem 7.2 (Margulis-Russo). *Let E be monotone. Then $d\mathbf{P}(E)/dp = \mathbf{I}^p(E)$.*

Proof. Temporarily, let $\mathbf{I}_i^{p_1, \dots, p_n}(E)$ denote the chance that $X^i \neq X$, where the $X_j \in \{0, 1\}$ are chosen uniformly at random with $\mathbf{P}(X_j = 1) = p_j$. Define $\mathbf{I}^{p_1, \dots, p_n}(E) = \sum \mathbf{I}_i^{p_1, \dots, p_n}(E)$. It suffices to show $\partial \mathbf{P}(E)/\partial p_i = \mathbf{I}_i^{p_1, \dots, p_n}(E)$, from which we can use the chain rule. We can write E as the union of two disjoint events E_0 and E_1 , where $E_0 = E \cap \{\chi_E(X^i) \neq \chi_E(X)\}$, and $E_1 = E \cap \{\chi_E(X^i) = \chi_E(X)\}$. Now E_1 does not depend on the value of X_1 at all, so $\mathbf{P}(E_1)$ is independent of p_i , and so

$$\frac{\partial \mathbf{P}(E_1)}{\partial p_i} = 0$$

On the other hand, by monotonicity, E_0 then equals the probability that $\chi_E(X^i) \neq \chi_E(X)$, intersected with the event $X_i = 1$. These two events are independent, so $\mathbf{P}(E_0) = p_i \mathbf{P}(\chi_E(X^i) \neq \chi_E(X))$. The latter probability does not depend on the index i , so

$$\frac{\partial \mathbf{P}(E_0)}{\partial p_i} = \mathbf{P}(\chi_E(X^i) \neq \chi_E(X)) = \mathbf{I}_1^{p_1, \dots, p_n}(E)$$

This completes the proof. \square

To analyze the critical exponent, we rely on two results we will prove later on using Fourier analysis, which allow us to upper bound the influence by the variance of a function.

Theorem 7.3 (Bourgain, Kahn, Kalai). *For any f and p , there exists i such that $\mathbf{I}_i^p(f) \gtrsim \mathbf{V}_p(f)[\log n/n]$, and $\mathbf{I}^p(f) \gtrsim \mathbf{V}_p(f) \log(1/\max \mathbf{I}_i^p(f))$.*

We also rely on a fact that, for exponentially many boolean inputs, the probability that a monotone event happens jumps from being unlikely to being likely in a very small range of p values, i.e. of length approximately $1/\log n$. Think like the majority function. Once $p > 1/2$, the chance of a vote passing grows rapidly in p .

Theorem 7.4 (Friedgut, Kalai). *If A is a monotone event, whose influences are the same for each index, and for $p = p_0$, if $\mathbf{P}(A) > \varepsilon$, then for $p \geq p_0 + c \log(1/\varepsilon)/\log n$, $\mathbf{P}(A) > 1 - \varepsilon$.*

Proof. If all the influences are the same, we find the total variance is

$$\mathbf{I}^p(E) \gtrsim \mathbf{V}_p(\chi_E) \log n = \min(\mathbf{P}(E), 1 - \mathbf{P}(E)) \log n$$

Now the Margulis-Russo formula yields that if $\mathbf{P}(E) \leq 1/2$,

$$\frac{d(\log \mathbf{P}(E))}{dp} = \frac{\mathbf{I}^p(E)}{\mathbf{P}(E)} \gtrsim \log n$$

Thus if $p \geq p_0 + c/\log n$, $\mathbf{P}(E) \geq 1/2$. Now if $\mathbf{P}(E) \geq 1/2$, then

$$\frac{d(\log(1 - \mathbf{P}(E)))}{dp} = -\frac{\mathbf{I}^p(E)}{1 - \mathbf{P}(E)} \lesssim -\log n$$

In order to make $\mathbf{P}(E) \geq 1 - \varepsilon$, we need $\log(1 - \mathbf{P}(E)) \leq \log \varepsilon$. To move from $\log(1/2)$ to $\log \varepsilon$, we need $p \geq p_0 + c/\log n + c \log(1/\varepsilon)/\log n = p_0 + c \log(1/\varepsilon)/\log n$. \square

Theorem 7.5 (FKG). *Any two increasing events are positively correlated.*

We now try and prove the critical exponent for the square lattice is $1/2$. First, we show $\theta(1/2) = 0$. Consider the ‘square annulus’ between 4^n and $3 \cdot 4^n$, and let E_n be the event that there is a ring in this annulus. Because the edge sets involved in each event are disjoint, the events are independent. Furthermore, the probability that there is a cross on each side of the annulus is bounded below by some constant $c > 0$. Adding more edges doesn’t hurt the probability that an event happens, so they are monotonic, and the FKG inequality says that $\mathbf{P}(E_n) = \mathbf{P}(A \cap B \cap C \cap D) \geq \mathbf{P}(A)\mathbf{P}(B)\mathbf{P}(C)\mathbf{P}(D) \geq c^4$. Thus infinitely many occur almost surely, so $\theta(1/2) = 0$, and in fact, there is almost surely no infinite cluster anywhere.

To form a contradiction, we assume the critical point is $1/2 + \delta$ instead of $1/2$. Given n , we form a $2n \times n$ box, and we let J_n denote the event that there is a crossing in the box, i.e a path from left to right and an edge from bottom to top.

Lemma 7.6. *As $n \rightarrow \infty$, $\max \mathbf{I}_e^p(J_n) \rightarrow 0$, where the maximum is taken over $1/2 \leq p \leq 1/2 + \delta/2$ and all edges e .*

Proof. If e is pivotal on a given input for J_n , that means there is a path from a vertex adjacent to e to a vertex a distance $n/2$ away. Thus by translation invariance, $\mathbf{I}_e^p(J_n)$ is bounded by half the probability that there is a path of length $n/2$ from the origin, which is uniformly bounded for $p \leq 1/2 + \delta/2$. As $\theta(1/2 + \delta/2) = 0$, these values must tend to zero as $n \rightarrow \infty$. \square

Lemma 7.7. *For some n , and with $p = 1/2 + \delta/2$, $\mathbf{P}(J_n) \geq 0.98$.*

Proof. We have already seen that $\inf \mathbf{P}(J_n) > 0$ when $p = 1/2$. If $\mathbf{P}(J_n) < 0.98$ for all n , then BKS shows that $d\mathbf{P}(J_n)/dp$ must tend to ∞ uniformly for all $1/2 \leq p \leq 1/2 + \delta/2$ as $n \rightarrow \infty$, by setting $\varepsilon < 0.02$. And this means that the probability of J_n increases massively over the range of p from $1/2$ to $1/2 + \delta/2$. \square

Now we show this implies that we almost surely get an infinite cluster. If we can cross a $2n \times n$ box with probability $1 - \varepsilon$, crossing lines gives a probability of $1 - 5\varepsilon$ of cross a $4n \times n$ box. Thus the probability that we cross a $4n$ by $2n$ box is greater than the probability that we cross the top and bottom of the graph. Thus

$$\begin{aligned} \mathbf{P}(\text{Top} \cup \text{Bottom}) &= \mathbf{P}(\text{Top}) + \mathbf{P}(\text{Bottom}) - \mathbf{P}(\text{Top})\mathbf{P}(\text{Bottom}) \\ &\geq 2(1 - 5\varepsilon) - (1 - 5\varepsilon)^2 \geq 1 - 5\varepsilon^2 \end{aligned}$$

If ε is small, this error is REALLY small. Thus almost surely all but finitely many of the J_n occur. Thus we just have to put these lines together in a way which guarantees

7.3 Conformal Invariance

Brownian motion is the limit of a random walk, and in the plane, is conformally invariant, in the sense that the image of a path of Brownian motion under an analytic map looks like the path of Brownian motion, up to a change in the measurement of time. Thus a random walk is conformally invariant ‘in the limit’. In some sense, percolation should also look ‘asymptotically’ conformally invariant in the limit.

Let us describe what this principle should look like when discretized. We consider a conformal map ϕ from the unit disk to some other simply connected domain D fixing the origin, and with $\phi'(0) > 0$. Now for any δ , we can consider the lattice $\delta\mathbf{Z}^2$, restricted to the interior of the unit disk. If C_δ is the cluster around the origin in the interior. Similarly, we define C'_δ to be the cluster around the origin in $\delta\mathbf{Z}^2$ restricted to D . Now $\phi(C_\delta)$ and C'_δ don’t even lie on the same lattice, but as $\delta \rightarrow 0$, they should still asymptotically describe the same law on space.

The simplest precise statement of conformal invariance was proved by Smirnov in 2001. Scale the hexagonal percolation problem by δ at the critical percolation value $p = 1/2$. If four points A, B, C , and D are chosen on the boundary of D , then the probability that there is a path from points on the boundary of D between A and B to points on the boundary of D between C and D converges as $\delta \rightarrow 0$. Furthermore, this convergent value is invariant under conformal mappings. In the case where D is a sidelength one equilateral triangle, A, B , and C are the three corner points, and D is on the line between A and C with distance x from C , the probability converges to x . By conformal invariance, this gives a general way to calculate the limiting probability. On \mathbf{Z}^2 , even this statement is still open.

Chapter 8

High Dimensional Probability

Theorem 8.1. *Let f be L Lipschitz. Then*

$$\gamma(x \in \mathbf{R}^n : |f(x) - M| > t) \leq 2 \exp(-t^2/2L^2)$$

Proof. Let \mathbf{H} be a half space with Gaussian measure $1/2$, i.e. the standard upper half plane $\mathbf{H} = \{x : x_1 \leq 0\}$. Then $\gamma(H_\varepsilon^c) = \mathbf{P}(N(0,1) \geq \varepsilon) \leq e^{-\varepsilon^2/2}$. Thus $\gamma(H_\varepsilon) \geq 1 - e^{-\varepsilon^2/2}$.

Next, if $A = \{x : f(x) \geq M\}$, we obtain

$$\gamma(x : f(x) \geq M - L\varepsilon) \geq \gamma(A_\varepsilon) \geq \gamma(\mathbf{H}_\varepsilon) \geq 1 - e^{-\varepsilon^2/2}$$

□

Theorem 8.2 (Isoperimetric Inequality in Gaussian Space). *Among all measurable sets A in \mathbf{R}^n with the same Gaussian measure, half spaces minimize the measure $\gamma_n(A_\varepsilon)$.*

Remark. We can replace M by $\mathbf{E}f$ using sub Gaussian centering.

Theorem 8.3. *Let X be sub Gaussian. Then the centered version satisfies $\|X - \mathbf{E}X\|_{\psi_2} \leq 2\|X\|_{\psi_2}$.*

Proof. We find

$$\|X - \mathbf{E}\|_p \leq \|X\|_p + \|\mathbf{E}X\|_p$$

Now $\mathbf{E}\|X\| \leq \|X\|_p$.

□

The normal distribution concentrates on an annulus of radius \sqrt{d} and width $O(1)$. Also, it concentrates on half spaces $\{X_1 \geq c\}$.

Bibliography

- [1] Larry Wasserman, *All of Statistics*
- [2] Walter Rudin, *Real and Complex Analysis*